

What you don't know about containers (and why it matters)



Sean Scott

Oracle ACE Director

Data on Kubernetes
Community Ambassador

Managing Principal Consultant



{ ON : The Beach }

Oracle on Docker

Running Oracle Databases in Linux Containers

Oracle on Docker

Running Oracle Databases in
Linux Containers

—
Sean Scott

apress®

Download a free sample chapter:
<https://oraclesean.com>

What Are Containers?



Containers are like lightweight Virtual Machines, right?



Electricity is like a water hose

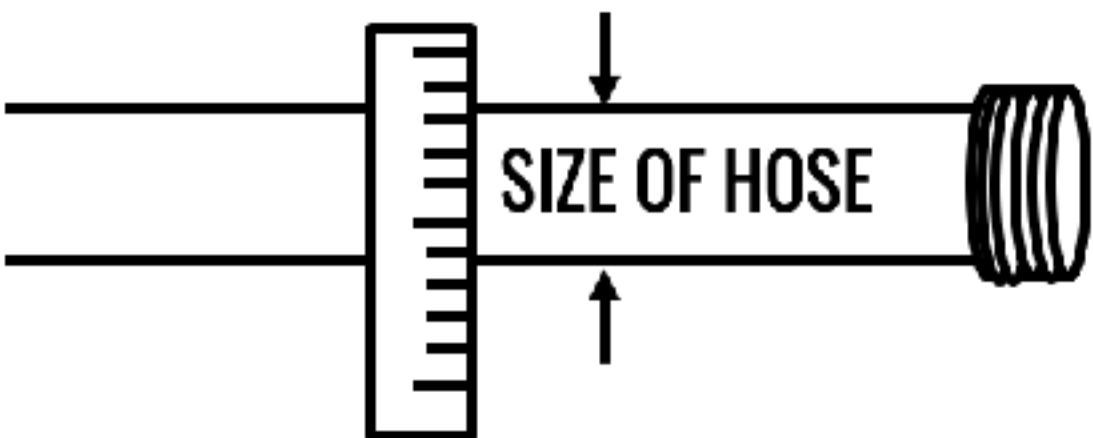
Voltage

Volts (V)



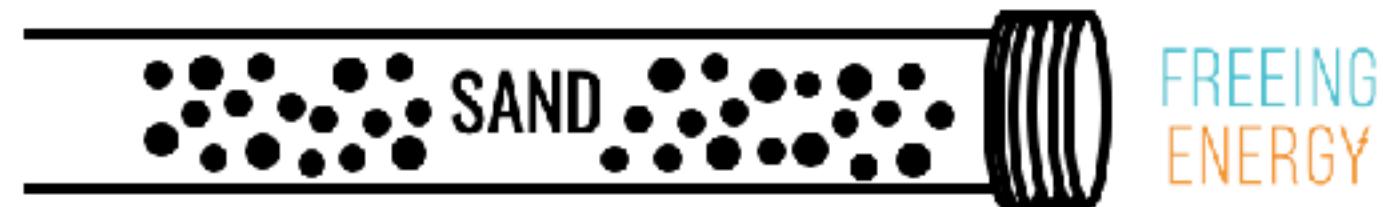
Current

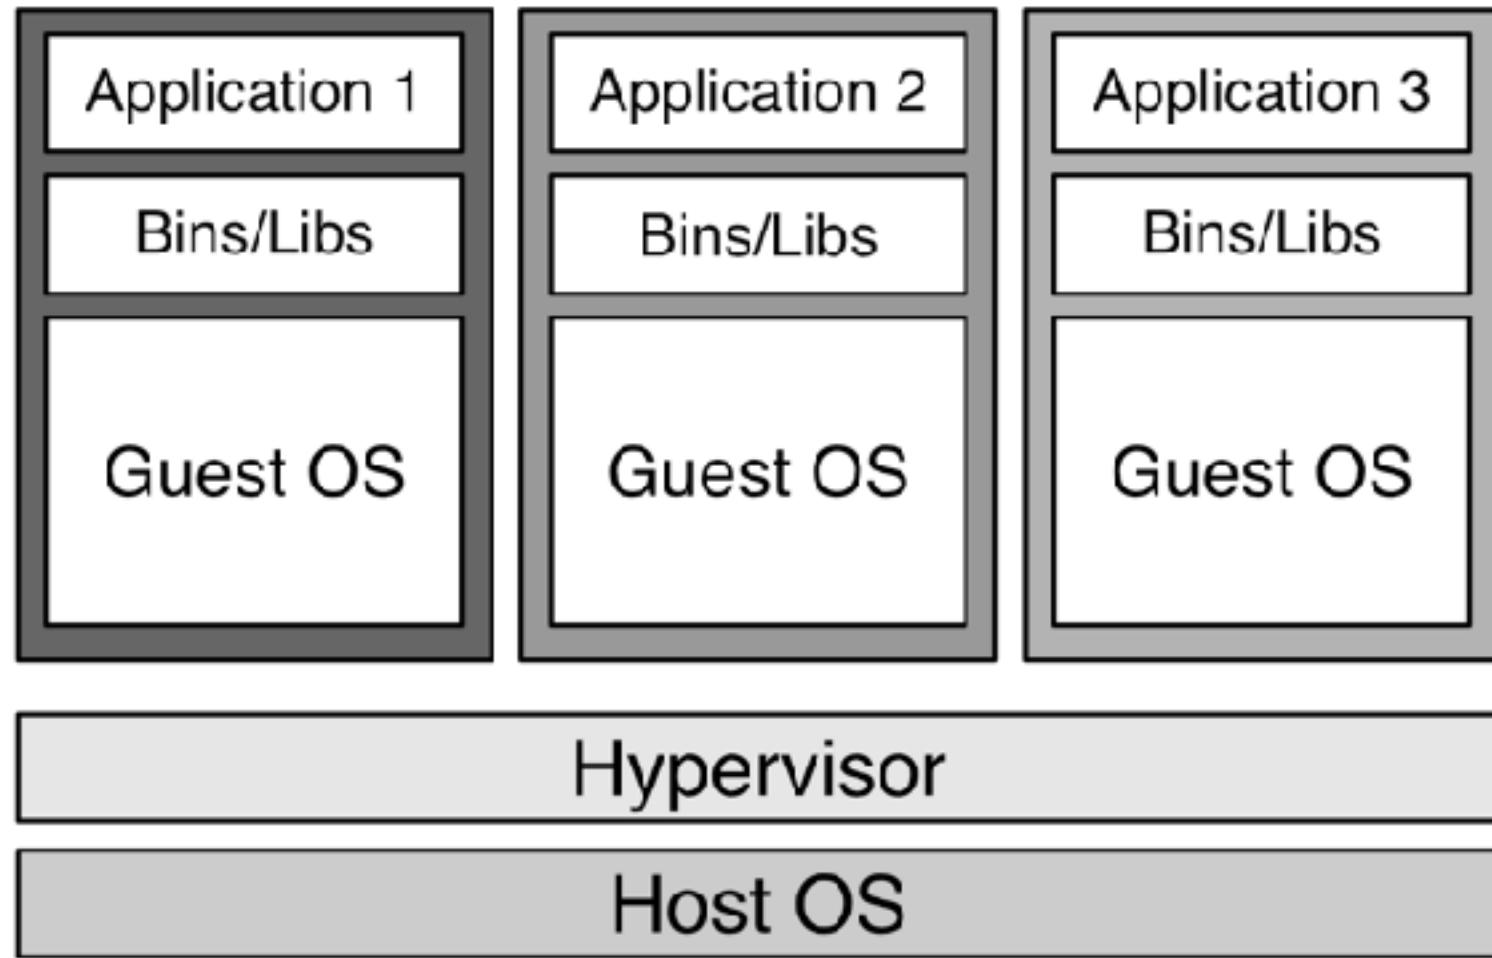
Amps (A or I)



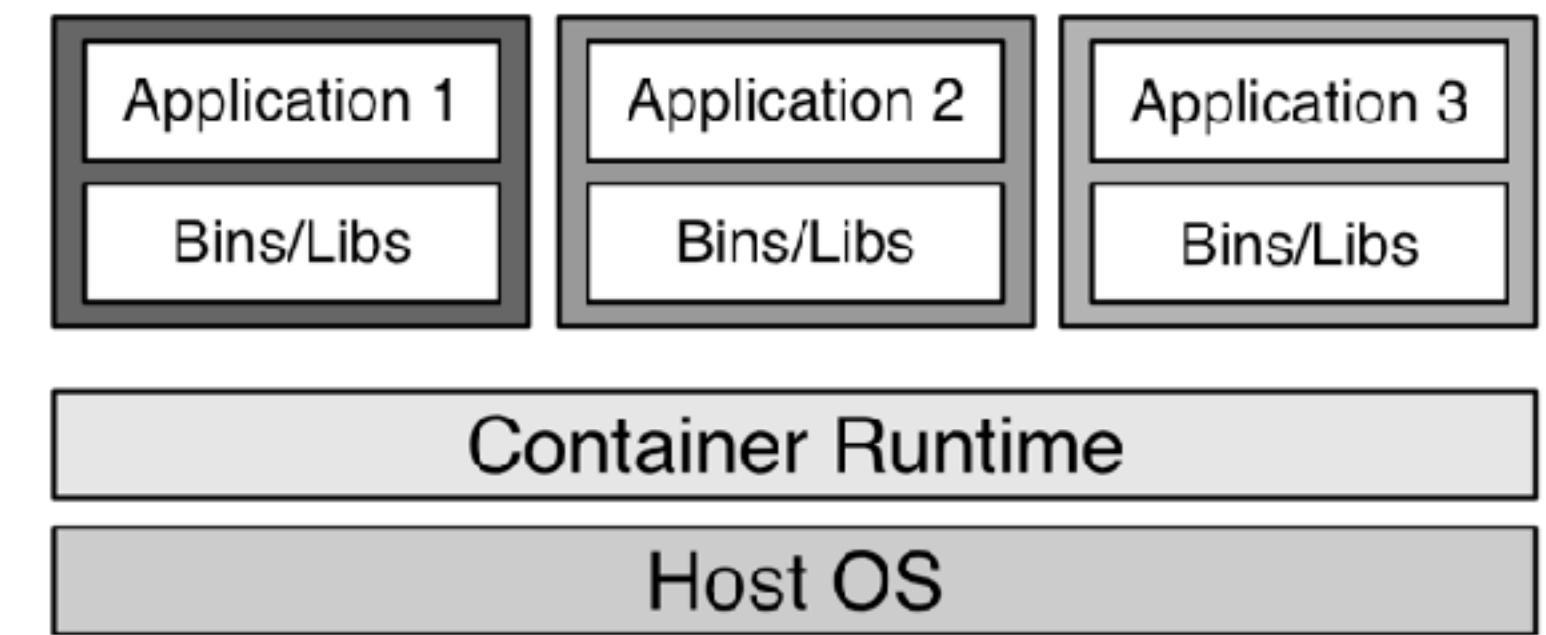
Resistance

Ohms (R or Ω)





Three full operating systems



Three application/executable directories

Virtual Machines: Bootable OS, often heavy

Containers: Support one application or service

Virtual Machines: Bootable OS, often heavy

A 100-page book of games & puzzles

Containers: Support one application or service

Virtual Machines: Bootable OS, often heavy

A 100-page book of games & puzzles

Containers: Support one application or service

A sheet of paper with a Tic-Tac-Toe grid

Games have structured playing
surfaces and rules.

Services deliver games to players.

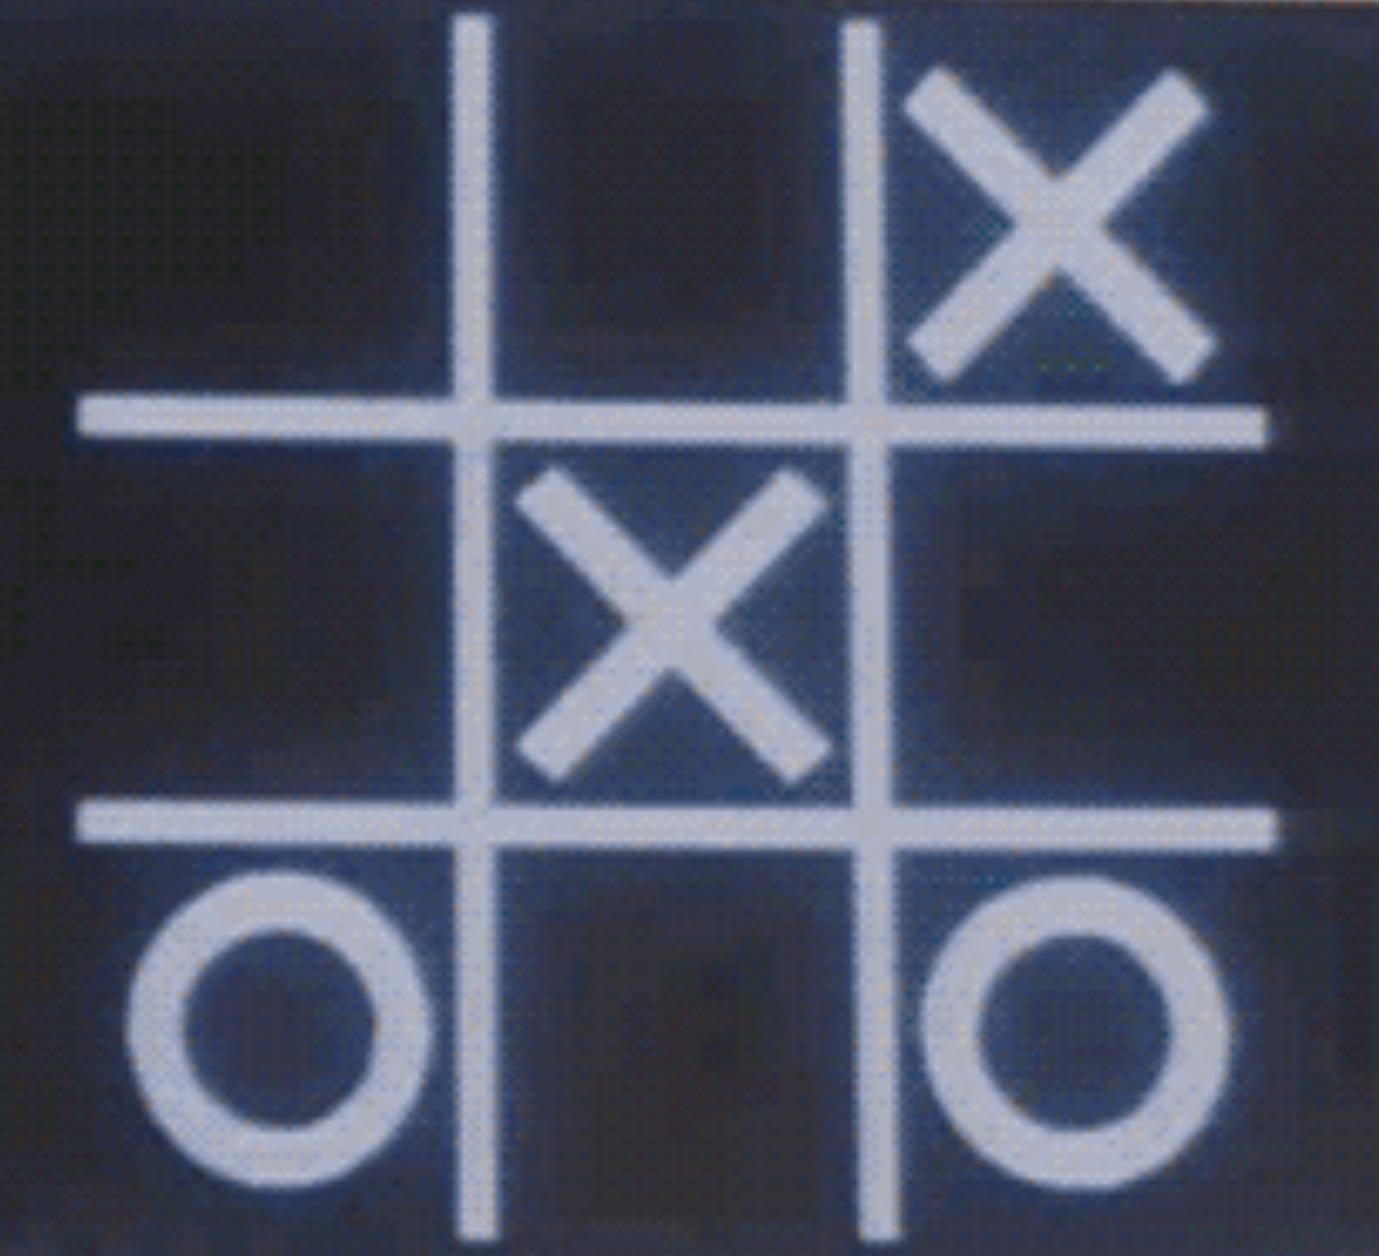
Images include the
game rules and playing surface.

Images are a minimal filesystem & metadata
for running a service.

Running an image starts a **Container**.

A container is:
**an isolated host process and
a Union Filesystem.**

SHALL WE PLAY A GAME?



UPD
0377

UPD
3652

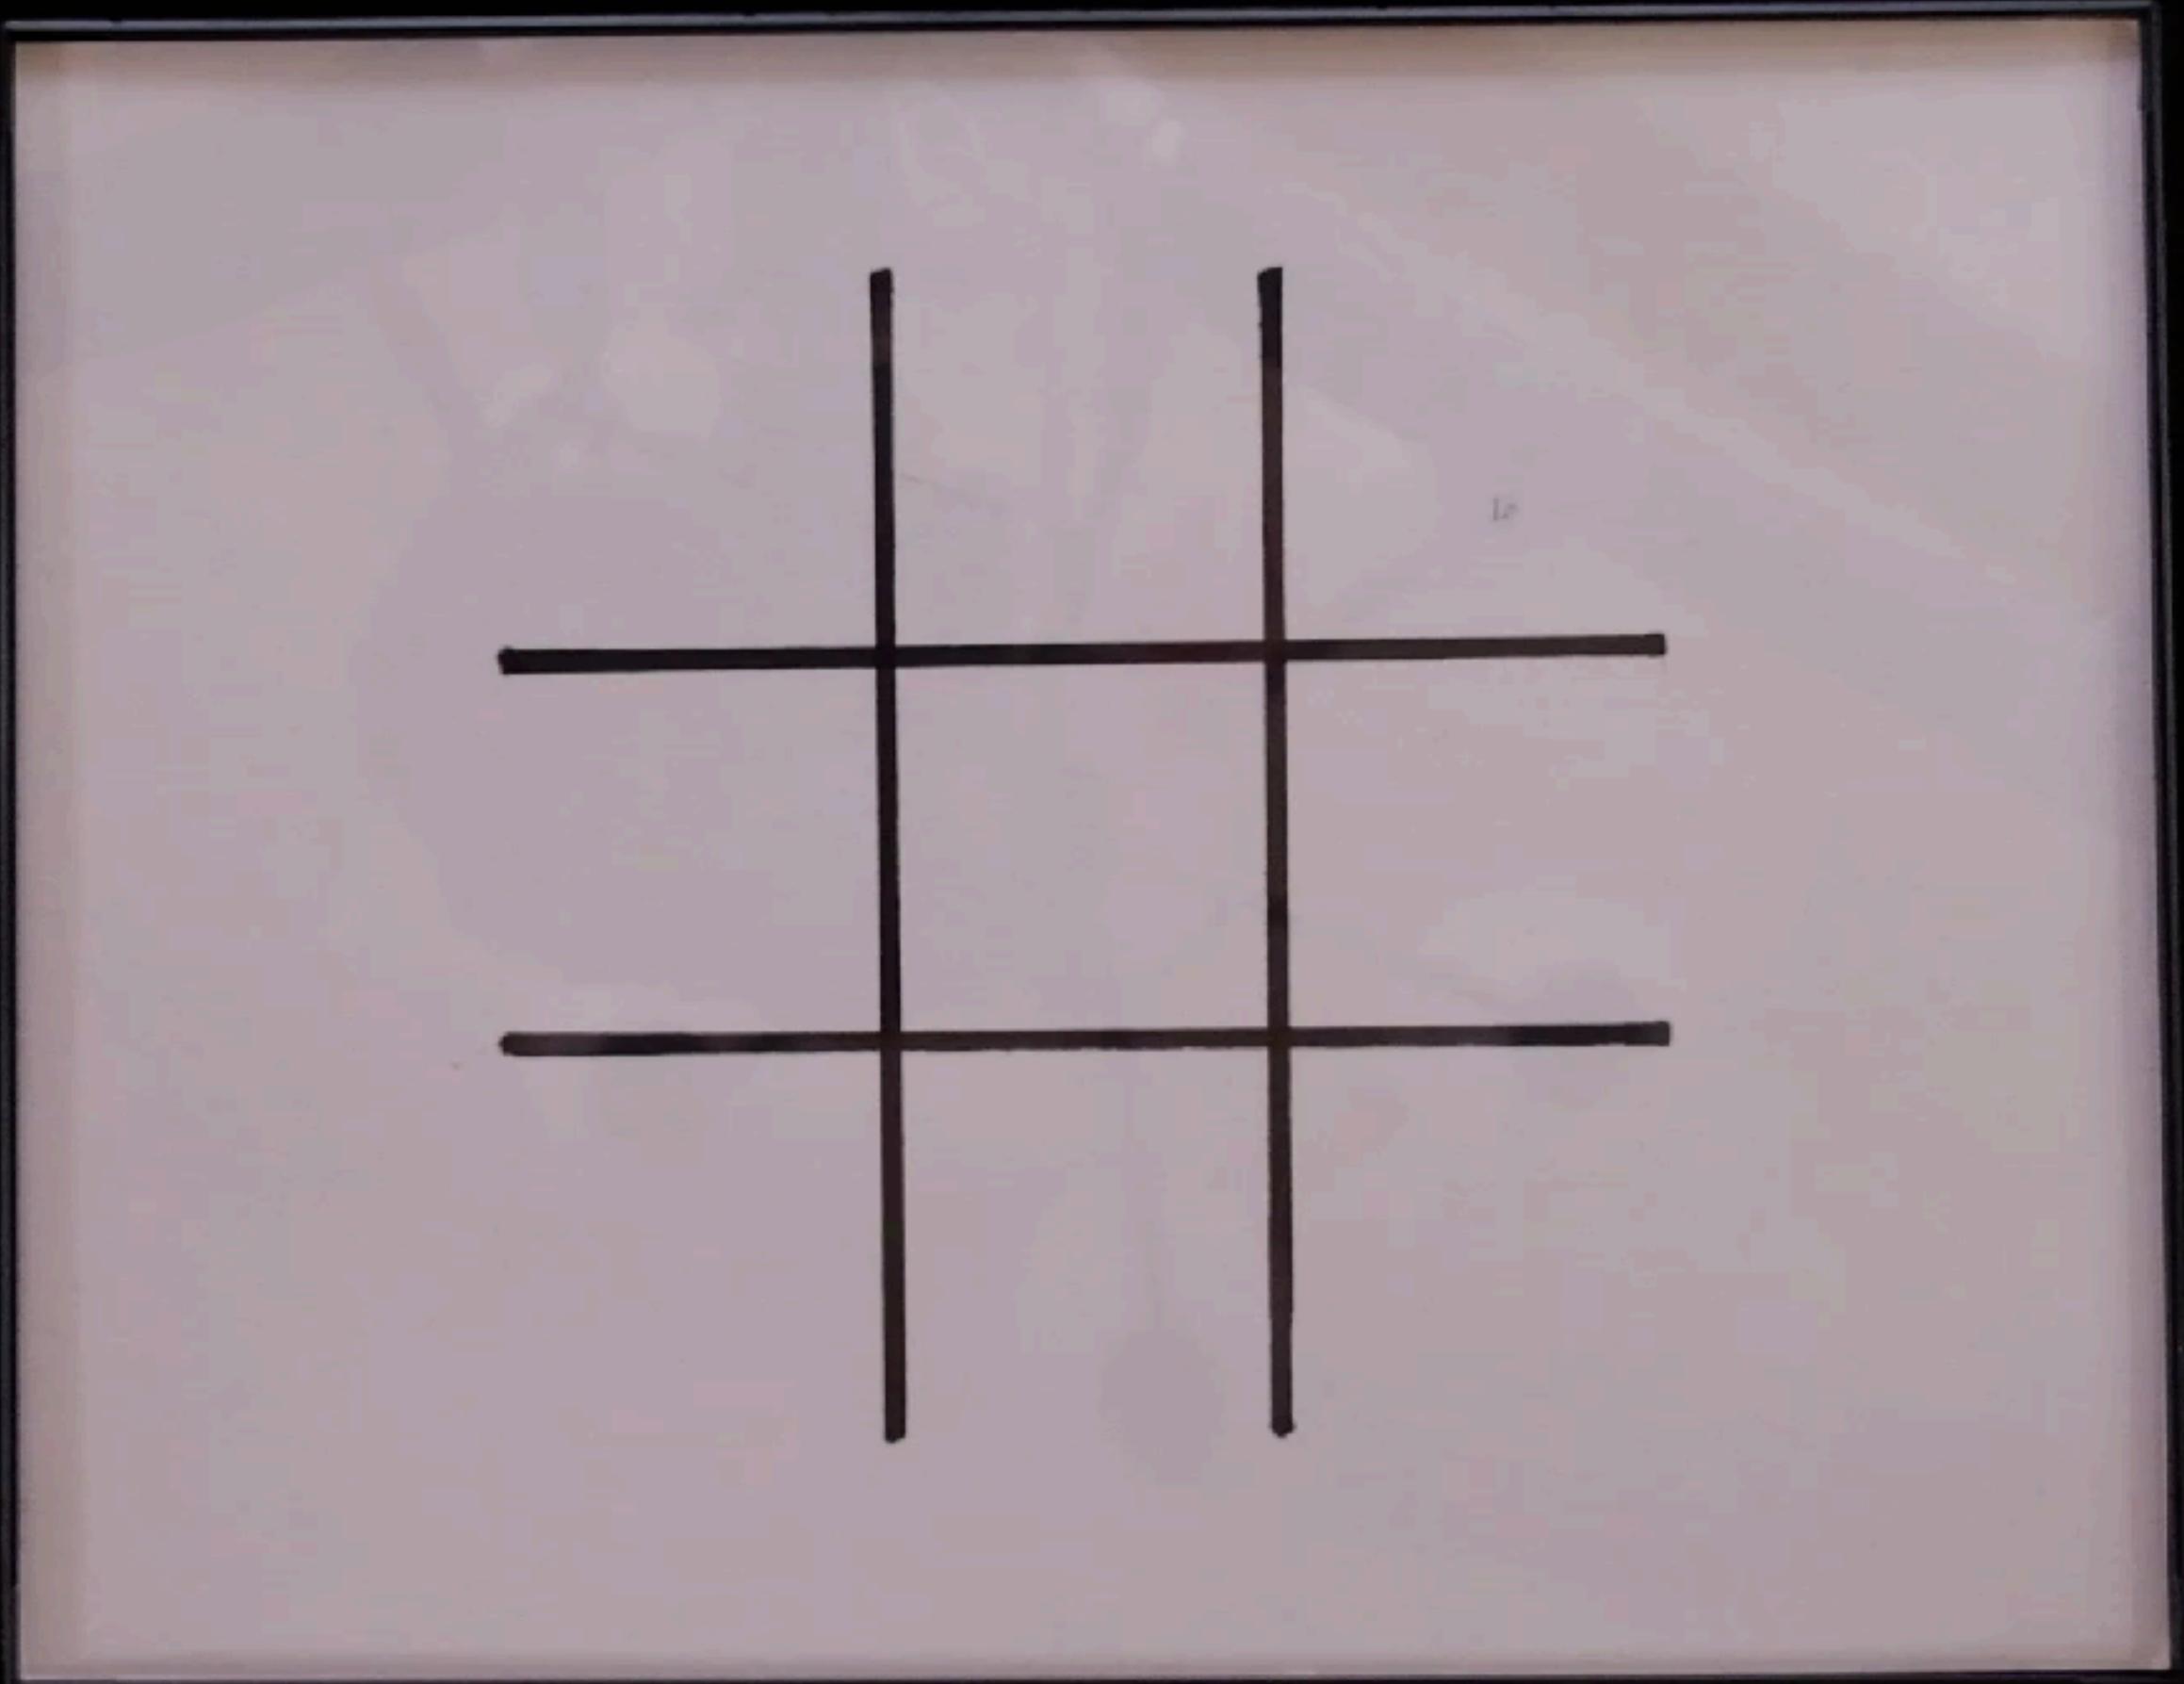
UPD
9080

UPD
4512

UPD
8884

UPD
6668





A **Union Filesystem** has three layers:

A Merge or Union Layer.

An Upper Layer.

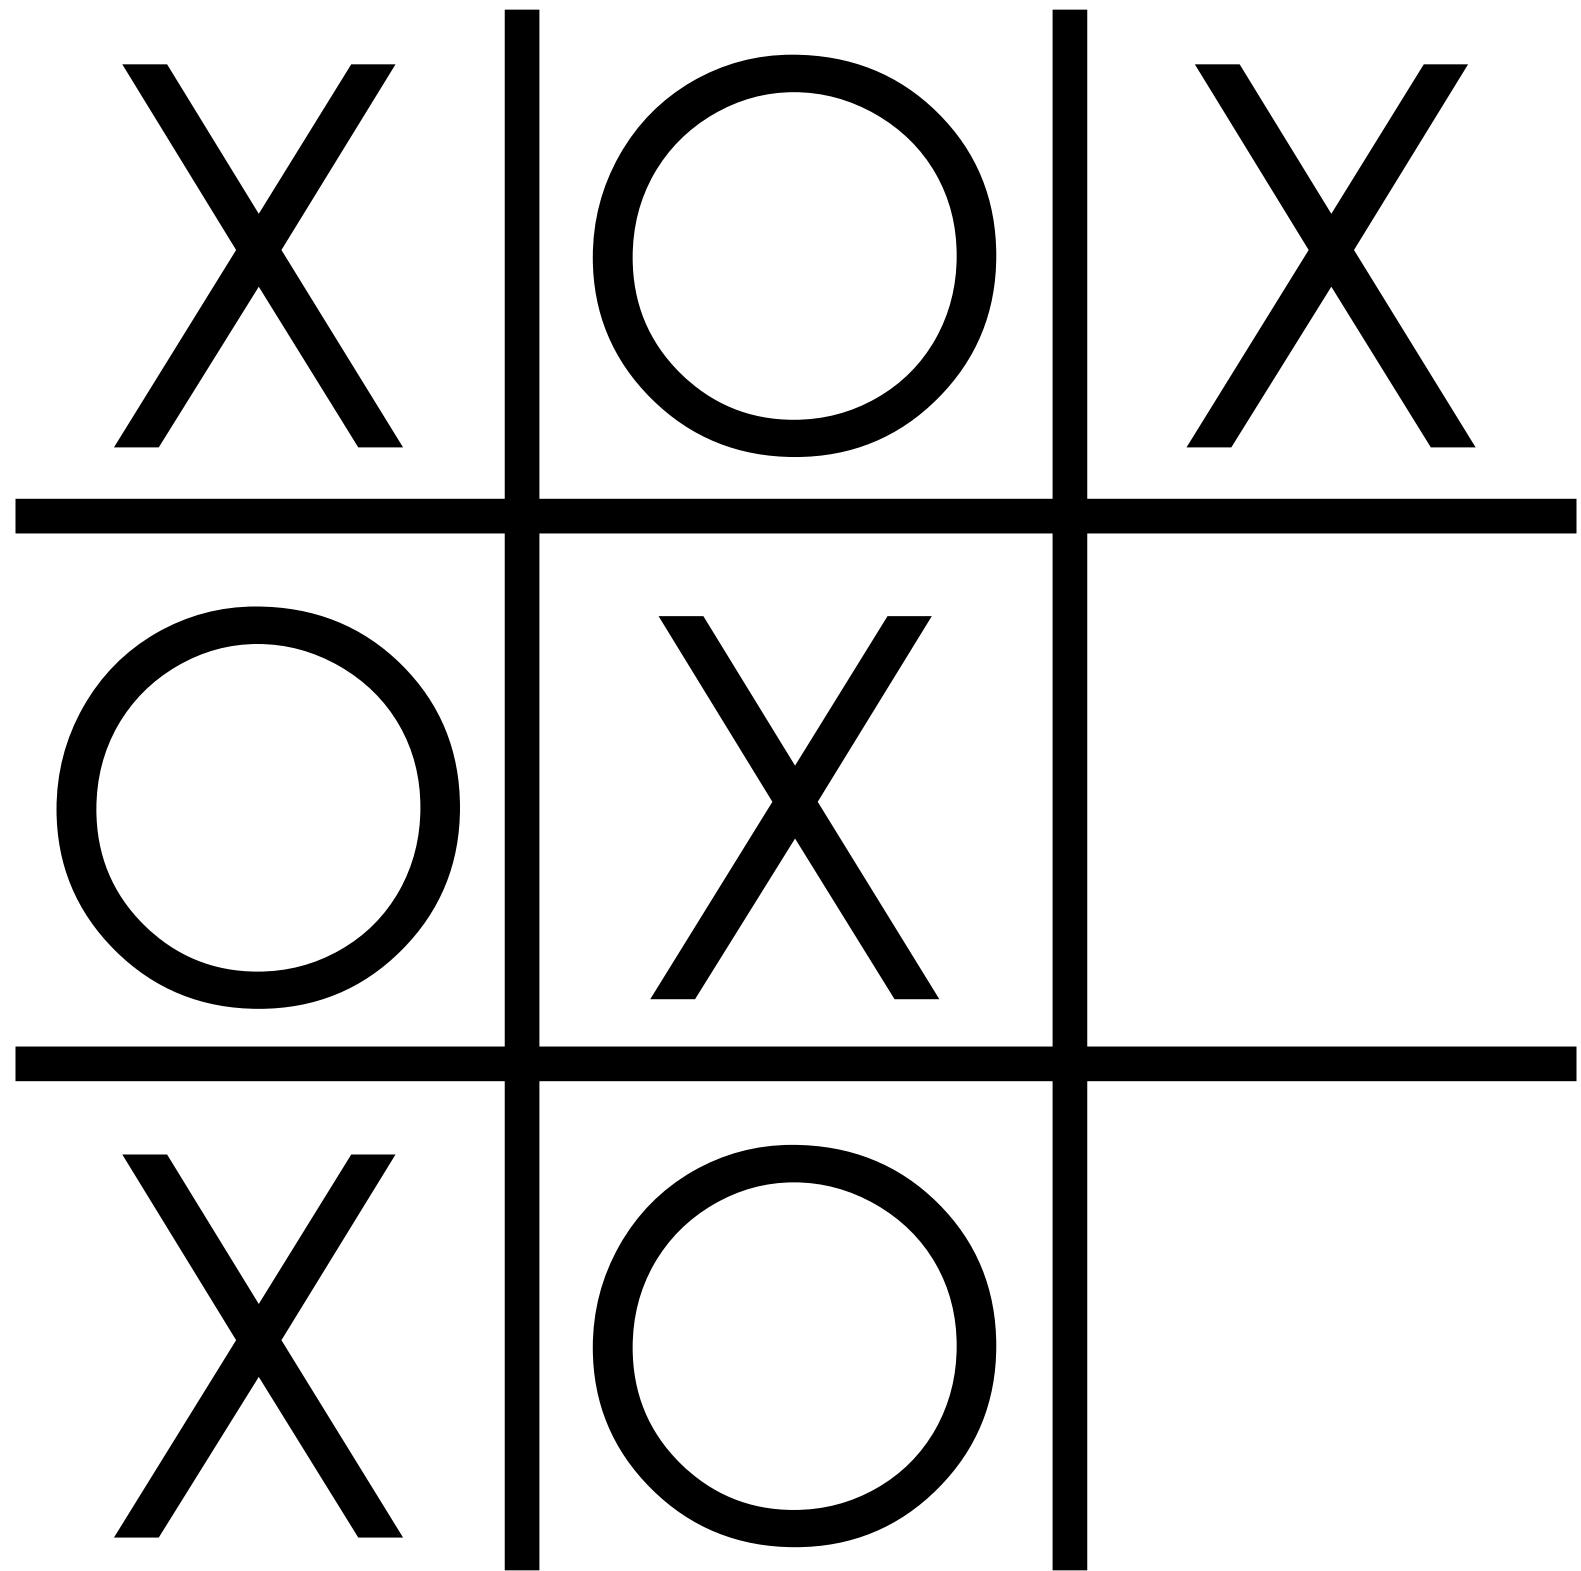
A Lower Layer.

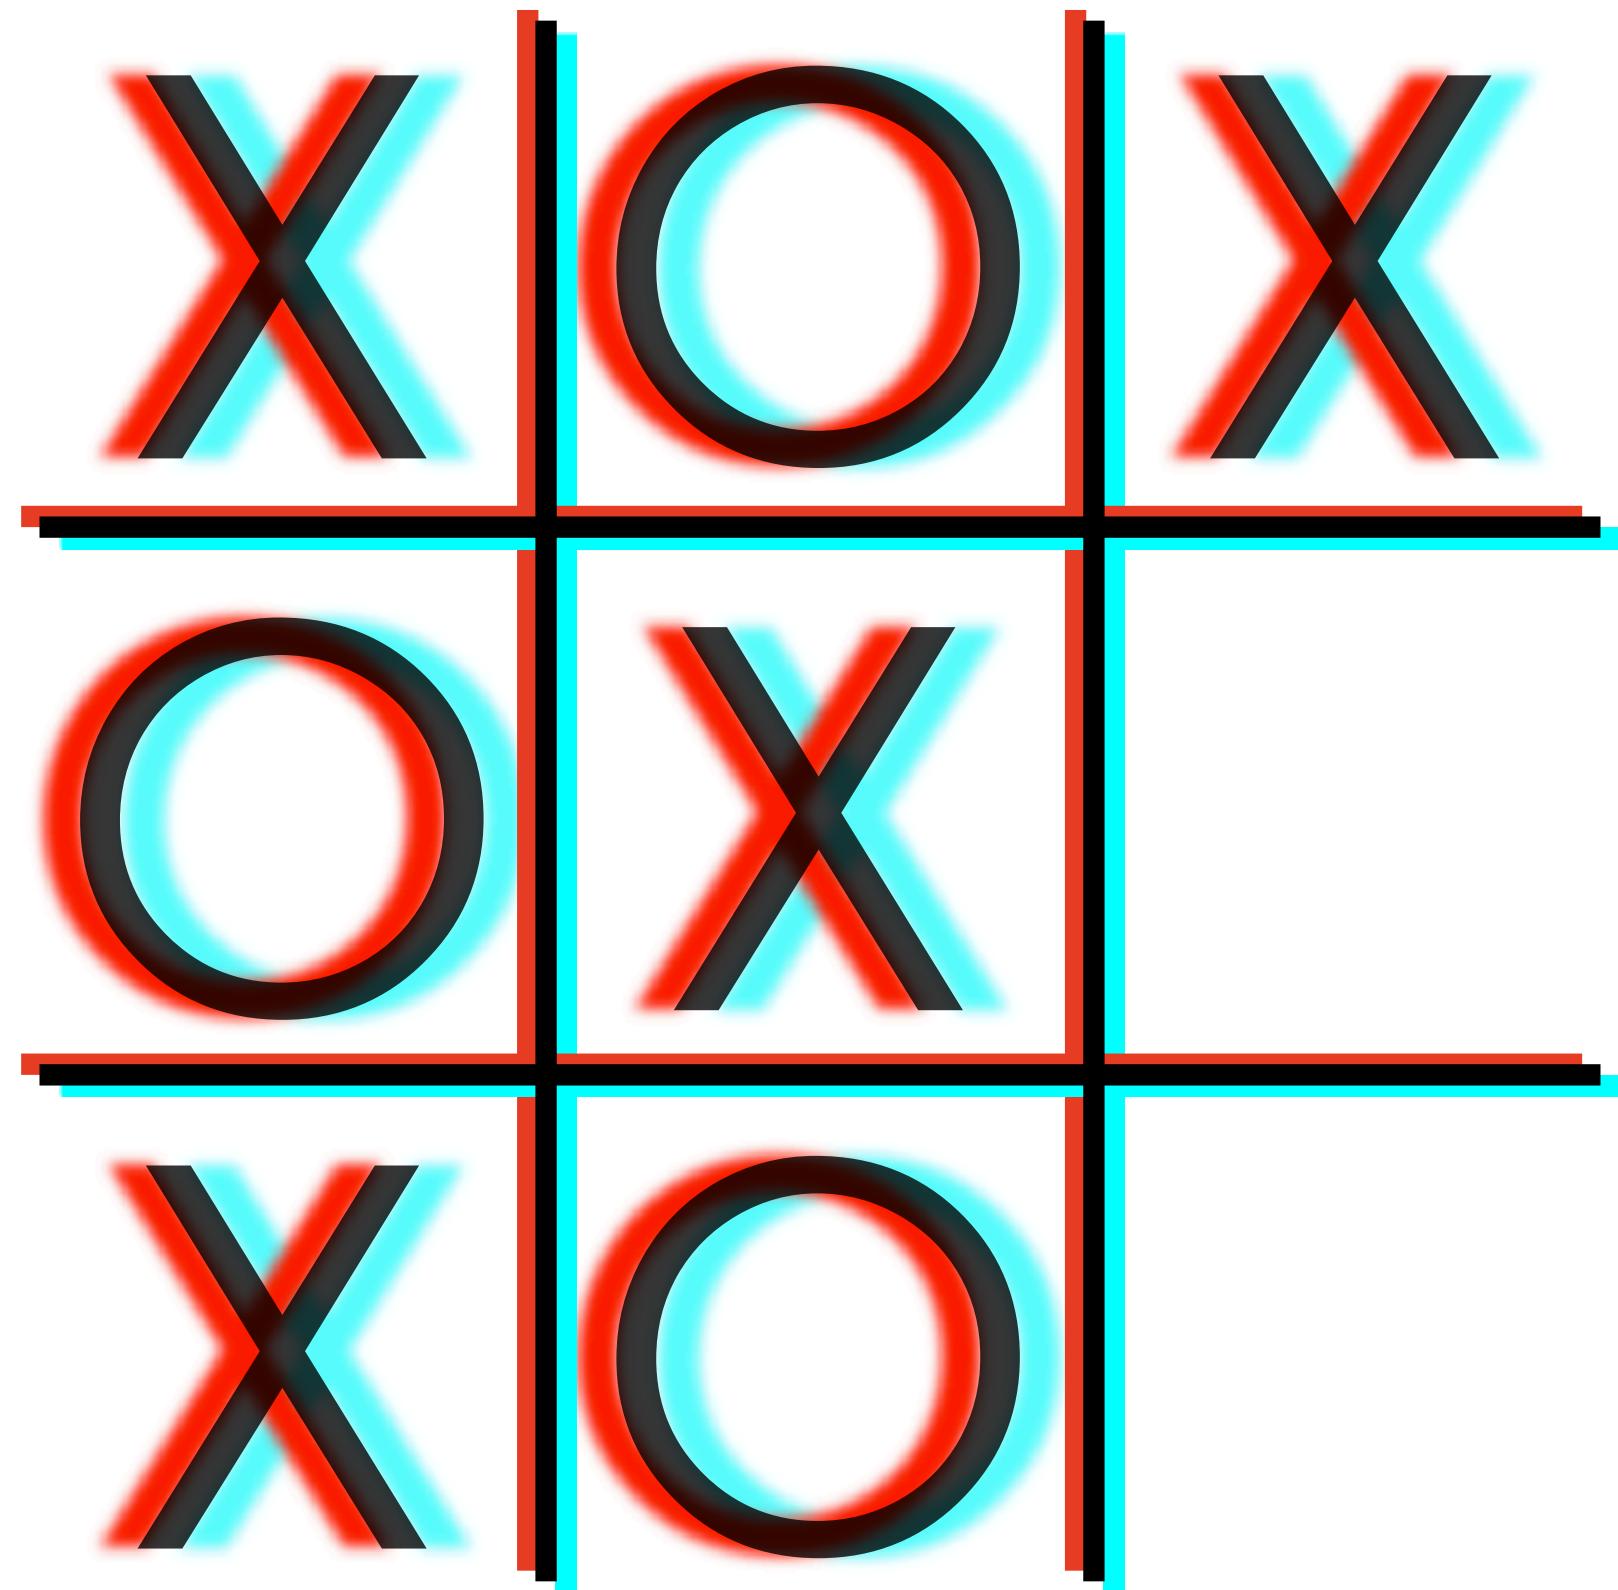
A **Union Filesystem** has three layers:

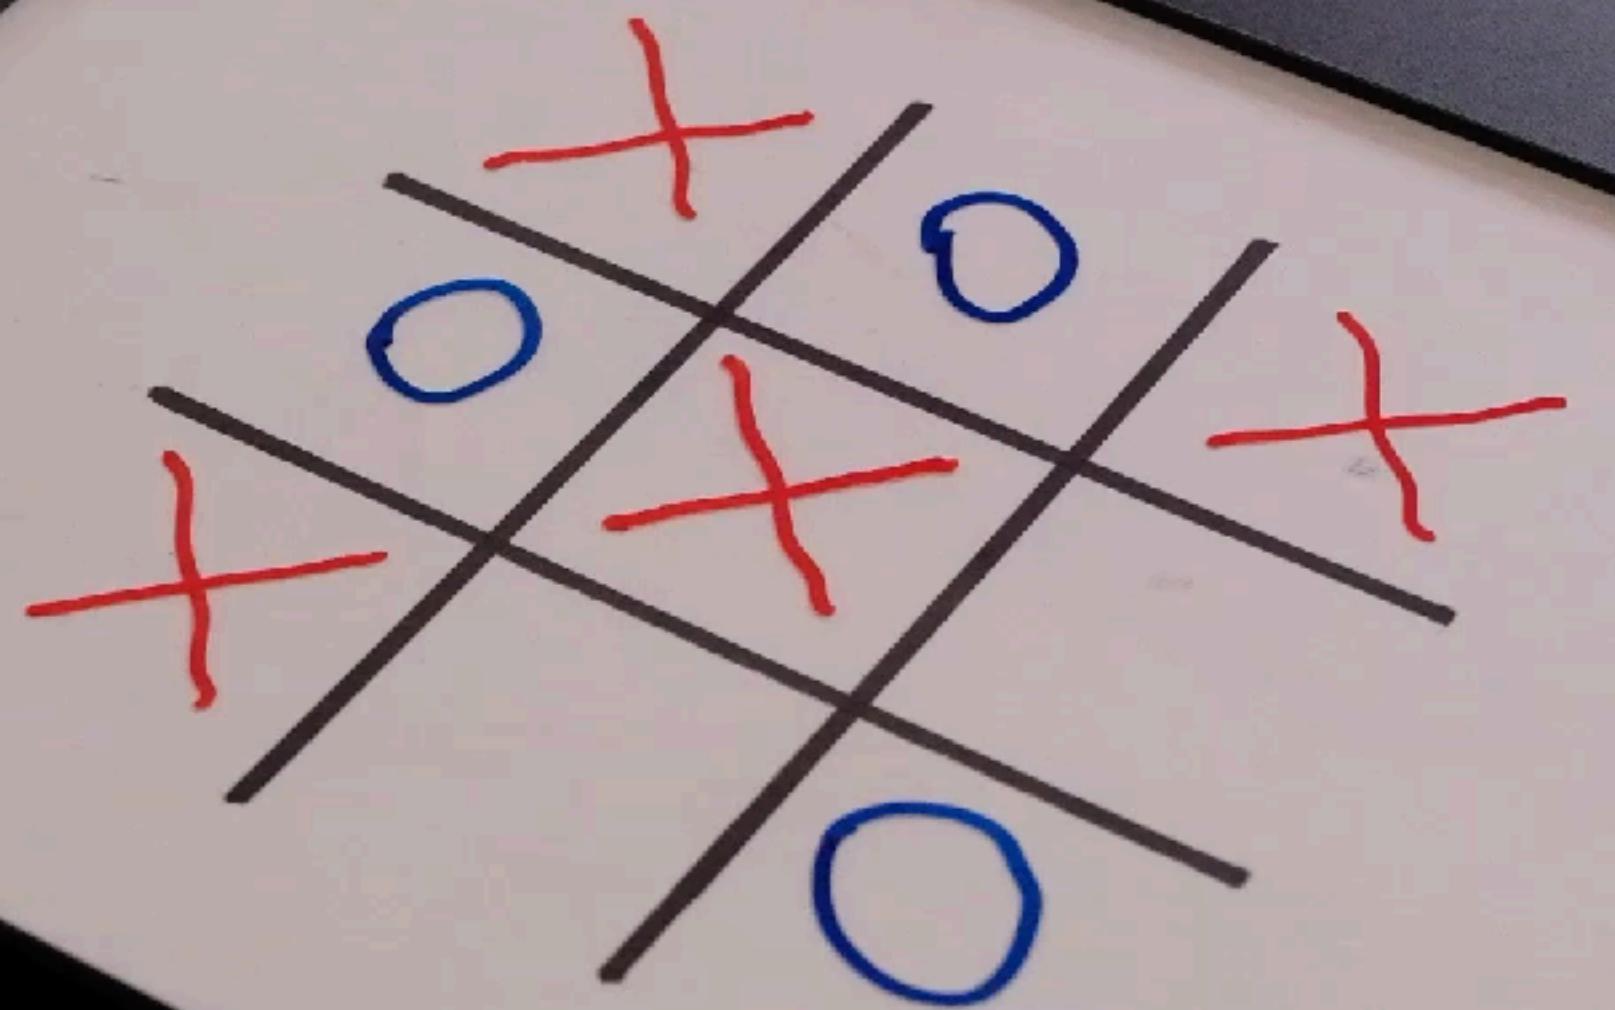
A Merge or Union Layer: The **game**.

An Upper Layer: The **player's moves**.

A Lower Layer: The **board and rules**.







The **Upper** (merge) layer isn't "real."

The Upper (merge) layer isn't "real."

It's a perceived (calculated) view
of two overlapping layers.

The **Upper** (merge) layer isn't "real."

It's a perceived (calculated) view
of **moves** and **the board**.

Union Filesystem Lab



```
# Create some directories:
```

```
mkdir -p ~/jotb/{image,c1_container,c1_work,container1}
```

```
# cd into the main directory:
```

```
cd ~/jotb
```

```
# Set the architecture:  
  
arch="$(arch)"  
arch="${arch/arm64/aarch64}"  
arch="${arch/amd64/x86_64}"  
  
# wget Alpine:  
  
wget -qO- https://dl-cdn.alpinelinux.org/alpine/v3.19/releases/"${arch}"/  
alpine-minirootfs-3.19.1-"${arch}".tar.gz \  
| tar xvz -C ~/jotb/image
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=c1_container, \
workdir=c1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=c1_container, \
workdir=c1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=c1_container, \
workdir=c1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=c1_container, \
workdir=c1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \  
-o lowerdir=image, \  
upperdir=c1_container, \  
workdir=c1_work \  
overlay container1
```

```
# List files in different layers:
```

```
ls -l ~/jotb/image
```

```
ls -l ~/jotb/c1_container
```

```
ls -l ~/jotb/container1
```

```
# Create files in different layers:
```

```
echo "AAA" > ~/jotb/image/AAA
```

```
echo "BBB" > ~/jotb/c1_container/BBB
```

```
# Change a file in the container layer:  
echo "aaa" > ~/jotb/c1_container/AAA
```

```
# Check the contents of file AAA in each layer:
```

```
cat ~/jotb/image/AAA
```

```
cat ~/jotb/c1_container/AAA
```

```
cat ~/jotb/container1/AAA
```

```
# Update PATH (just in case!)
```

```
PATH=$PATH:/bin
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/jotb/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/jotb/container1 \  
/bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/jotb/container1 \  
/bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/jotb/container1 \  
/bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/jotb/container1 \  
/bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/jotb/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/jotb/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/jotb/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/jotb/container1 \  
/bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/jotb/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/jotb/container1 \  
    /bin/sh
```

Check the hostname:

hostname

```
# Change the hostname
```

```
hostname alpine  
hostname
```

```
# Explore the "container"
```

```
cat /etc/os-release
```

```
uname -a
```

```
whoami
```

```
echo $$
```

```
pwd
```

What files are here?

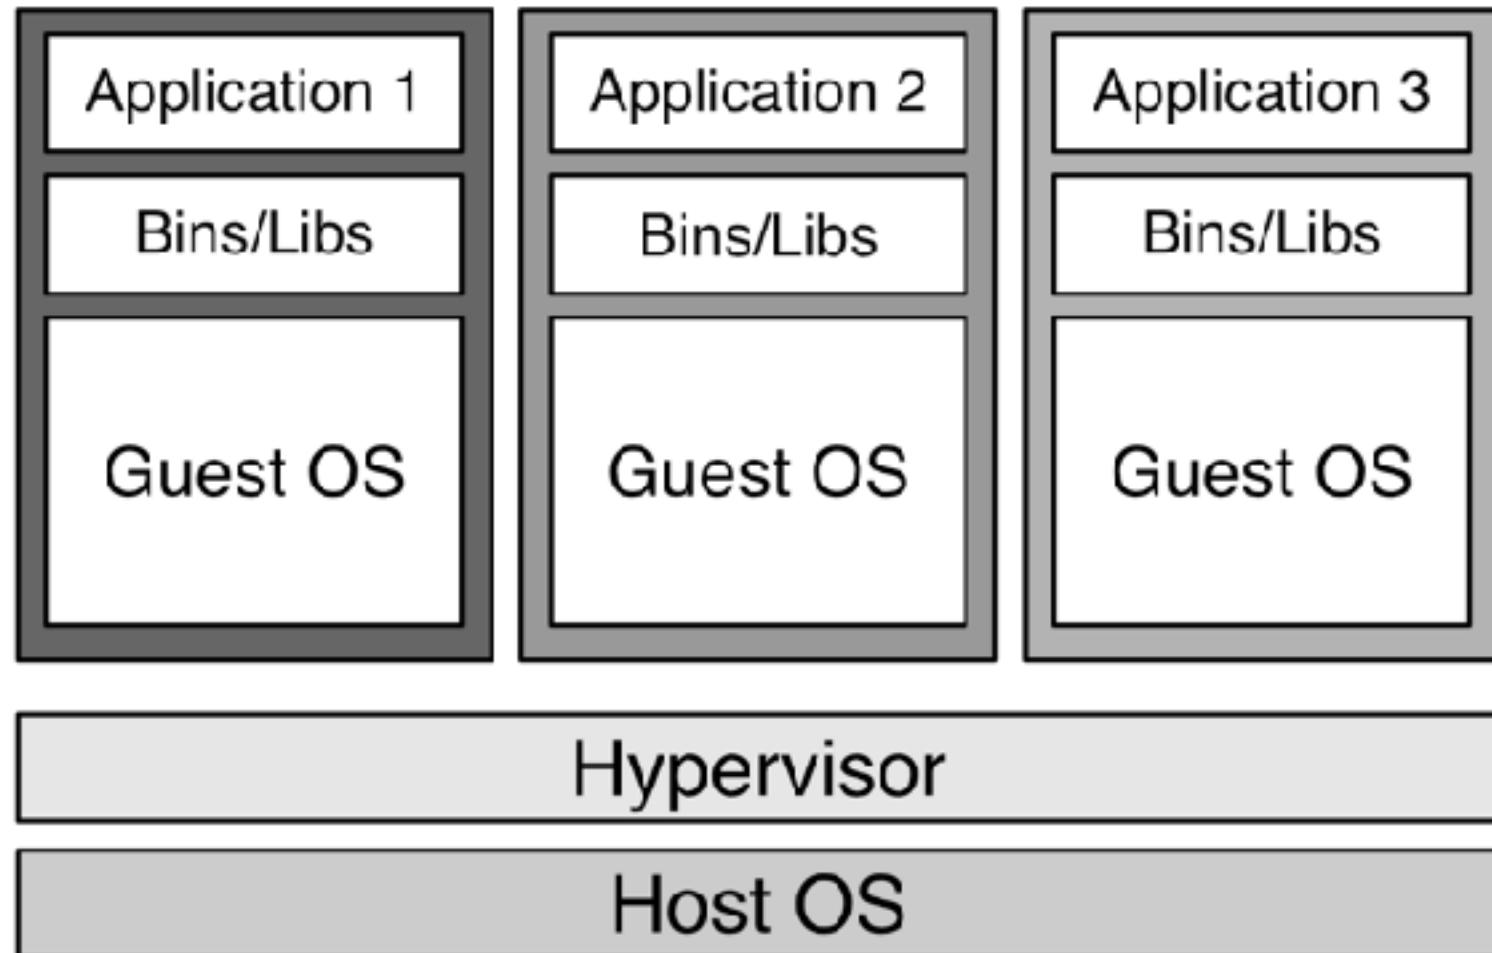
ls -l /

cat AAA

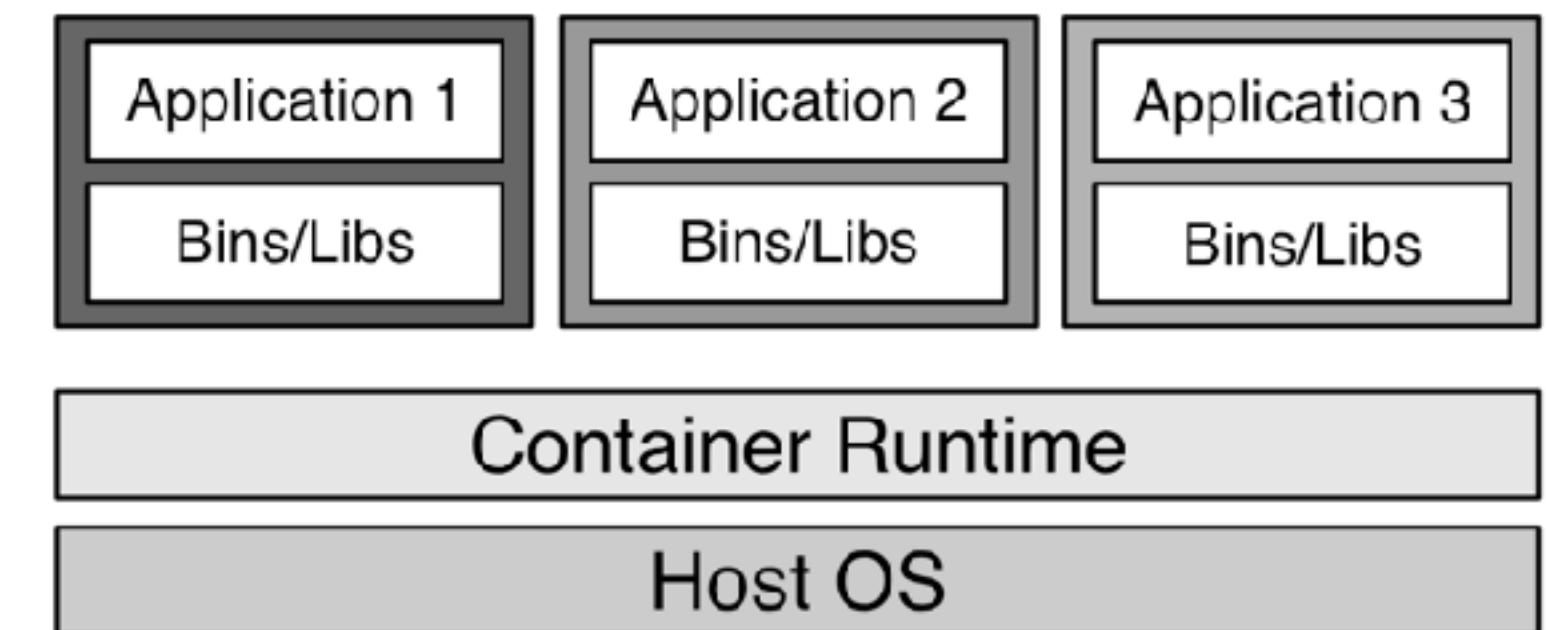
cat BBB

Containers are Magic!

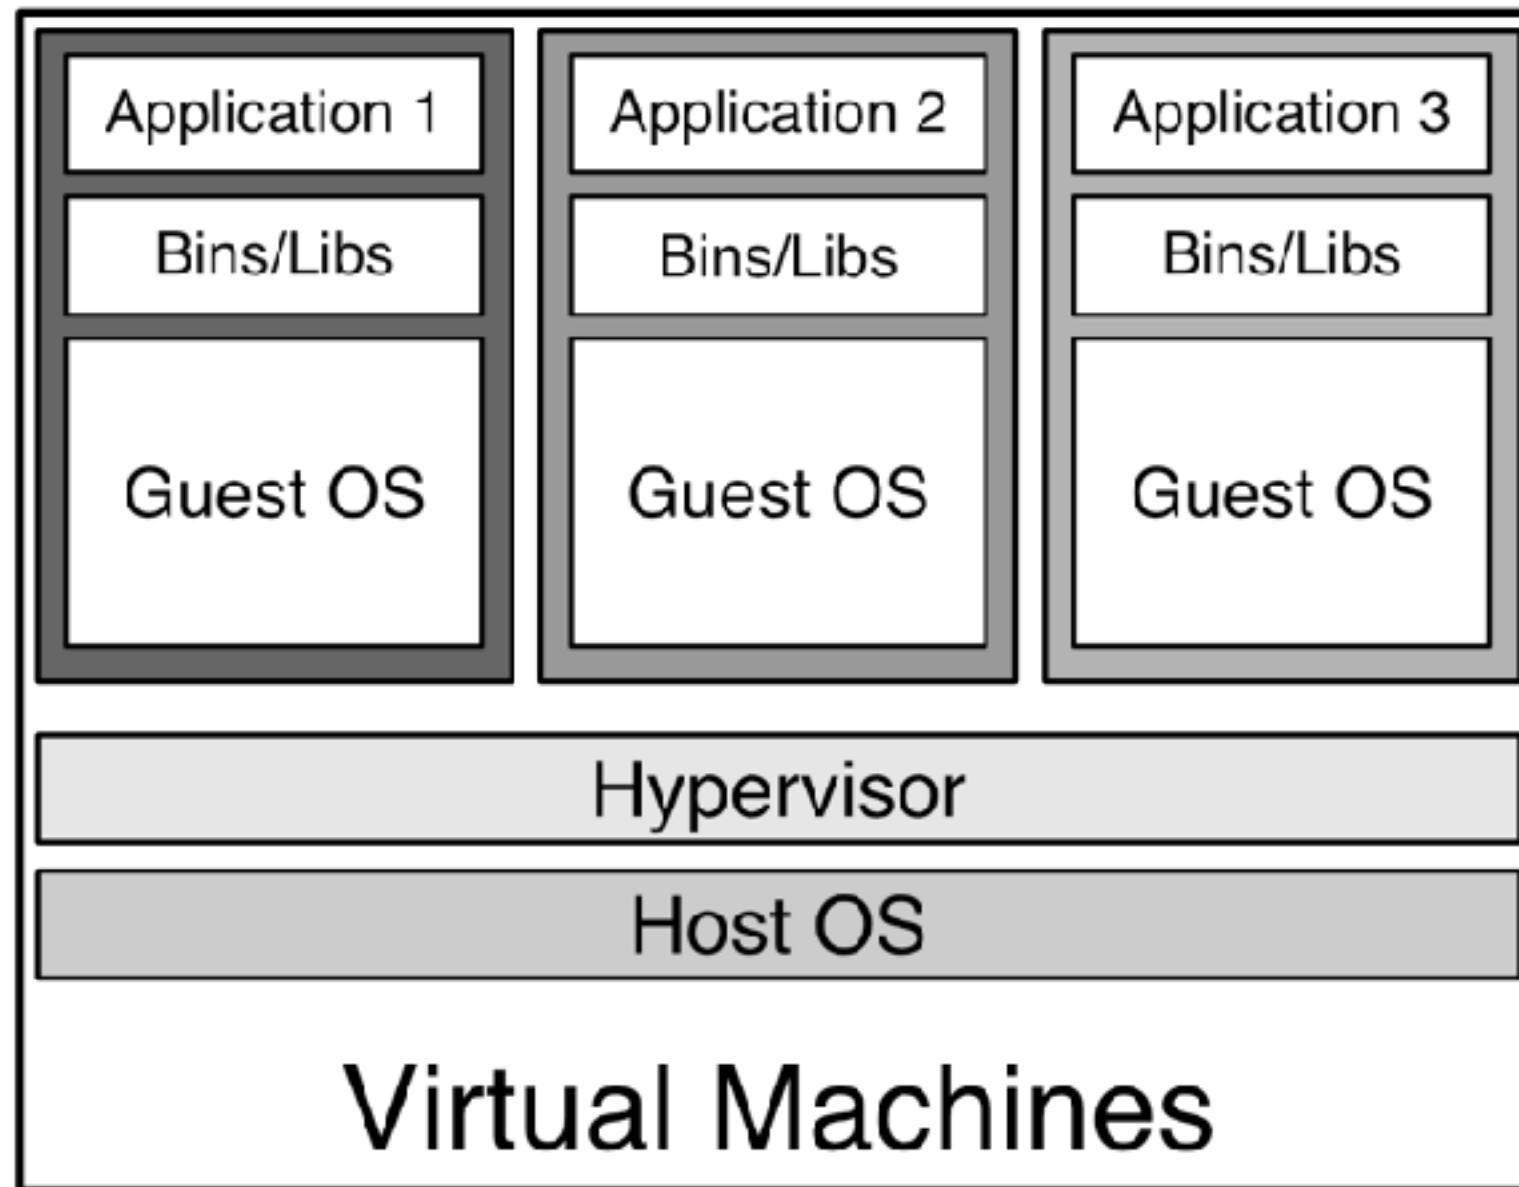




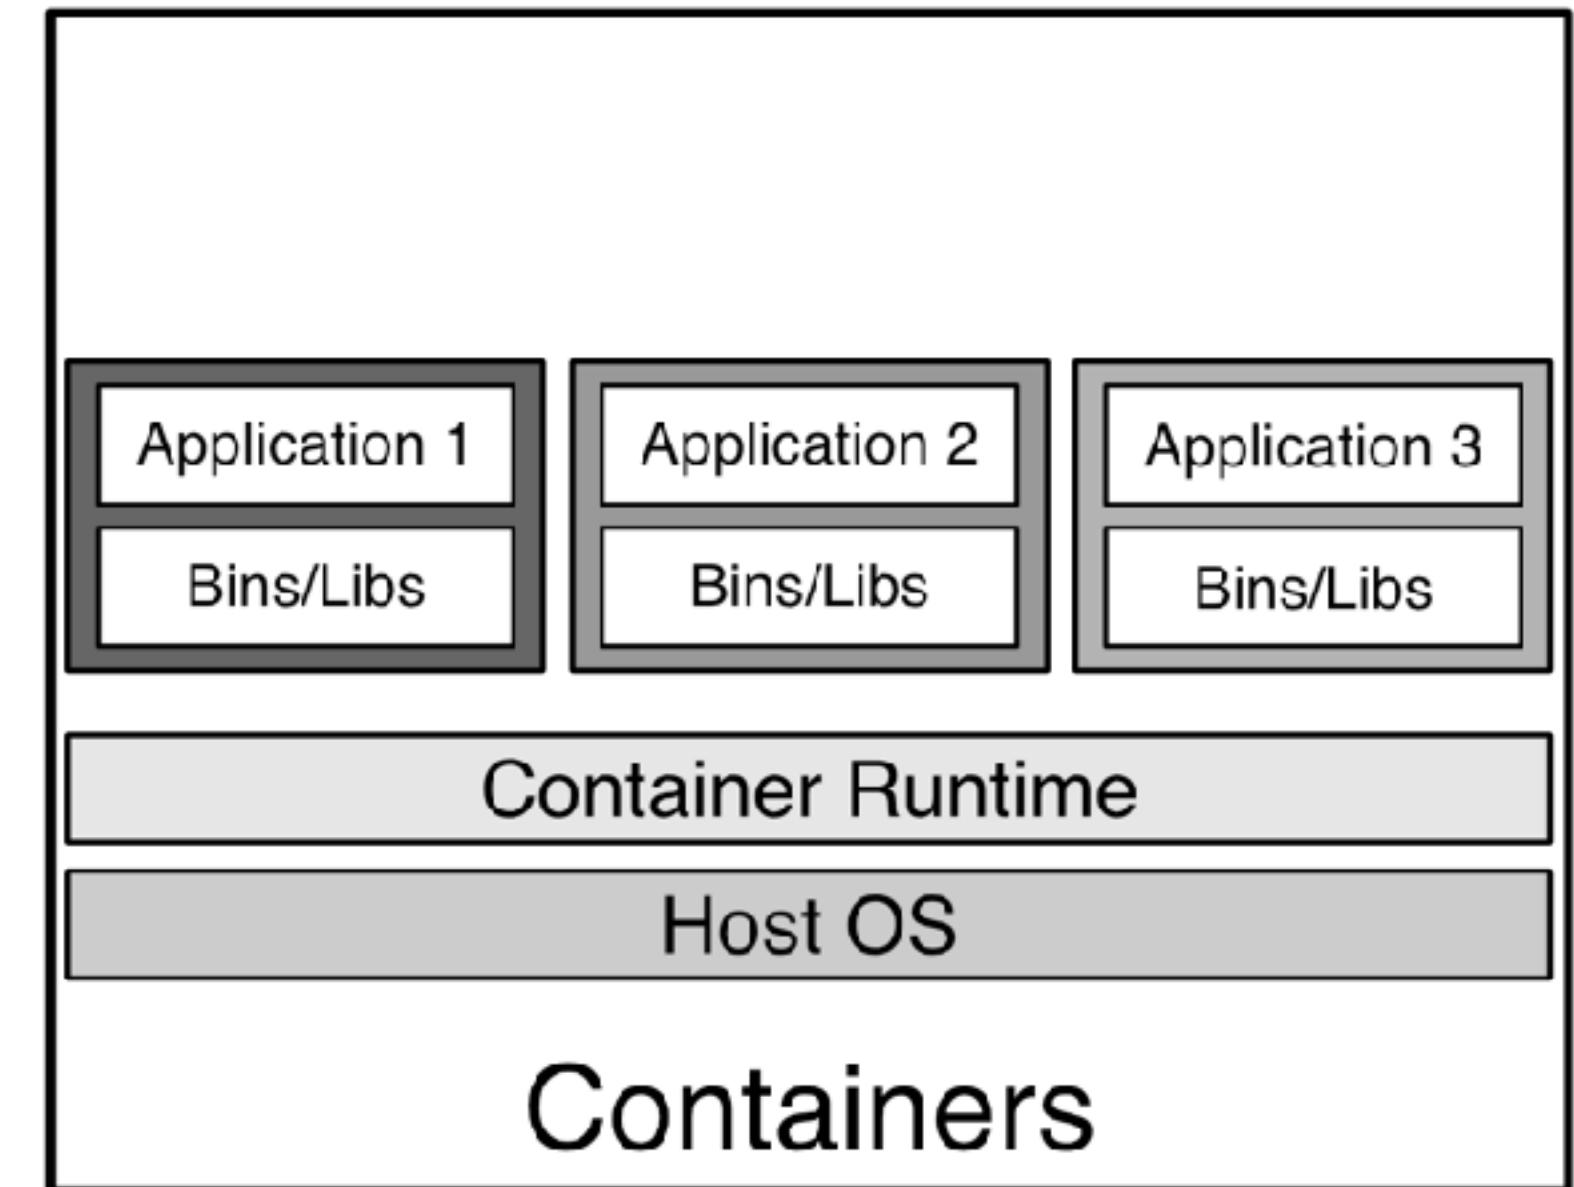
Three full operating systems



Three application/executable directories



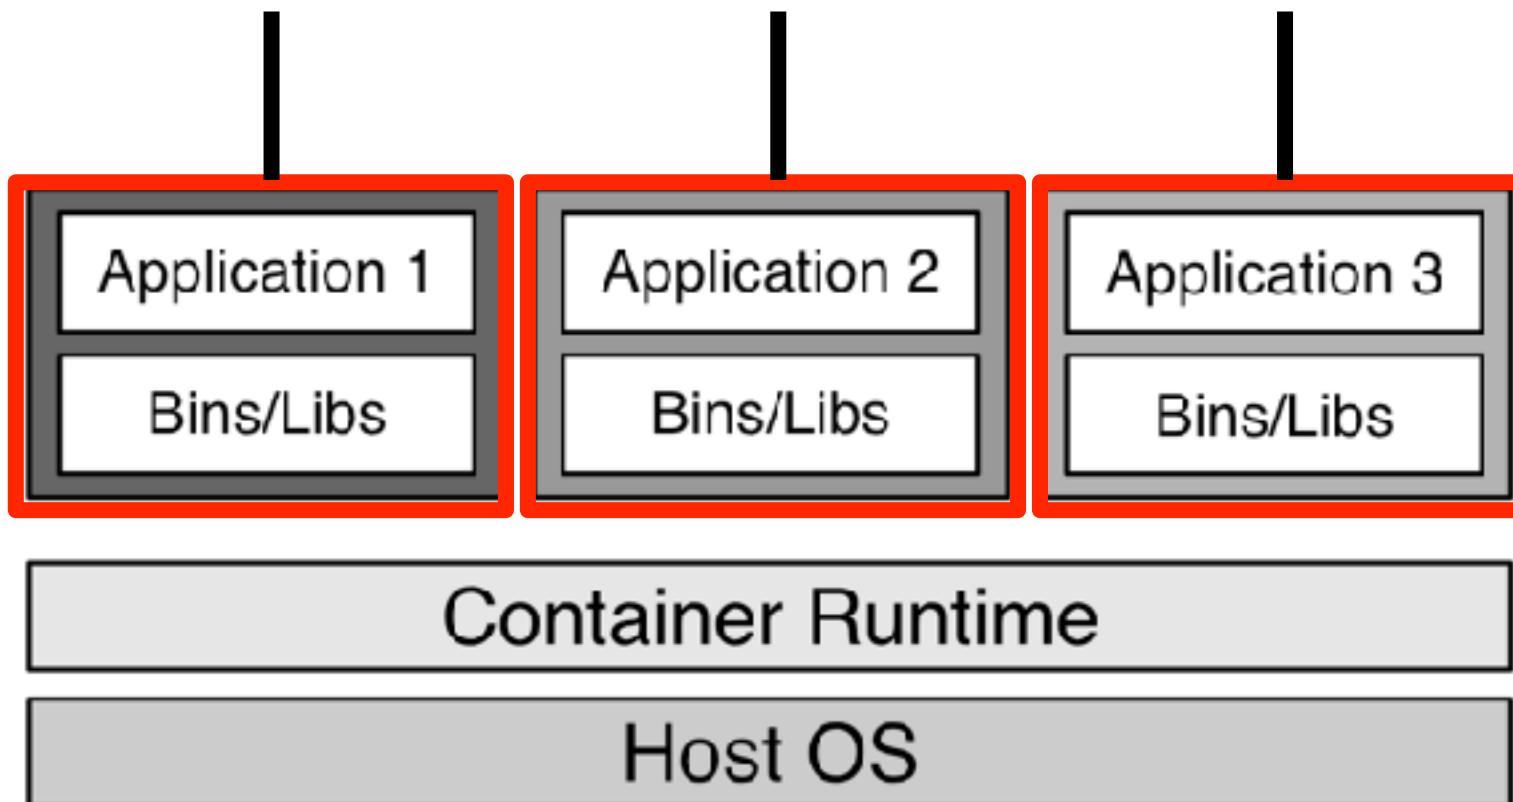
Three full operating systems



Three application/executable directories

How much space do three containers use?

500MB 500MB 500MB



$$\begin{array}{r} 500\text{MB} \\ + 500\text{MB} \\ + 500\text{MB} \\ \hline 1,500\text{MB} \end{array}$$

Multiple identical containers will share an **image**.

```
# Prepare directories for a second container  
mkdir -p ~/jotb/{c2_container,c2_work,container2}  
  
cd ~/jotb
```

```
# Create a second overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=c2_container, \
workdir=c2_work \
overlay container2
```

```
# Create a second overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=c2_container, \
workdir=c2_work \
overlay container2
```

```
# Create a second overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=c2_container, \
workdir=c2_work \
overlay container2
```

```
# Set the environment (just in case!)  
  
PATH=$PATH:/bin  
  
# "Run" the "container"  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/jotb/container2 \  
    /bin/sh
```

What files are here?

ls -l /

cat AAA

cat BBB

```
# Manipulate files within the container:
```

```
echo "111" > /AAA  
echo "222" > /BBB
```

```
# Check the contents of files from the host:
```

```
cat ~/jotb/image/AAA
```

```
cat ~/jotb/c2_container/AAA
```

```
cat ~/jotb/c2_container/BBB
```

Containers **persist** state (moves) in the run-time layer of its Union Filesystem.

Containers are **stateful**.

```
# Exit and restart the container:
```

```
exit
```

```
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/jotb/container2 \  
/bin/sh
```

```
# Confirm the file contents:
```

```
ls -l /
```

```
cat AAA
```

```
cat BBB
```

Container Properties

Dropping a **Container** removes its **Layers**.



```
# "Drop" the container:  
  
sudo umount container2  
  
rm -fr ~/jotb/container2  
rm -fr ~/jotb/c2_container  
rm -fr ~/jotb/c2_work
```

Containers are Ephemeral.

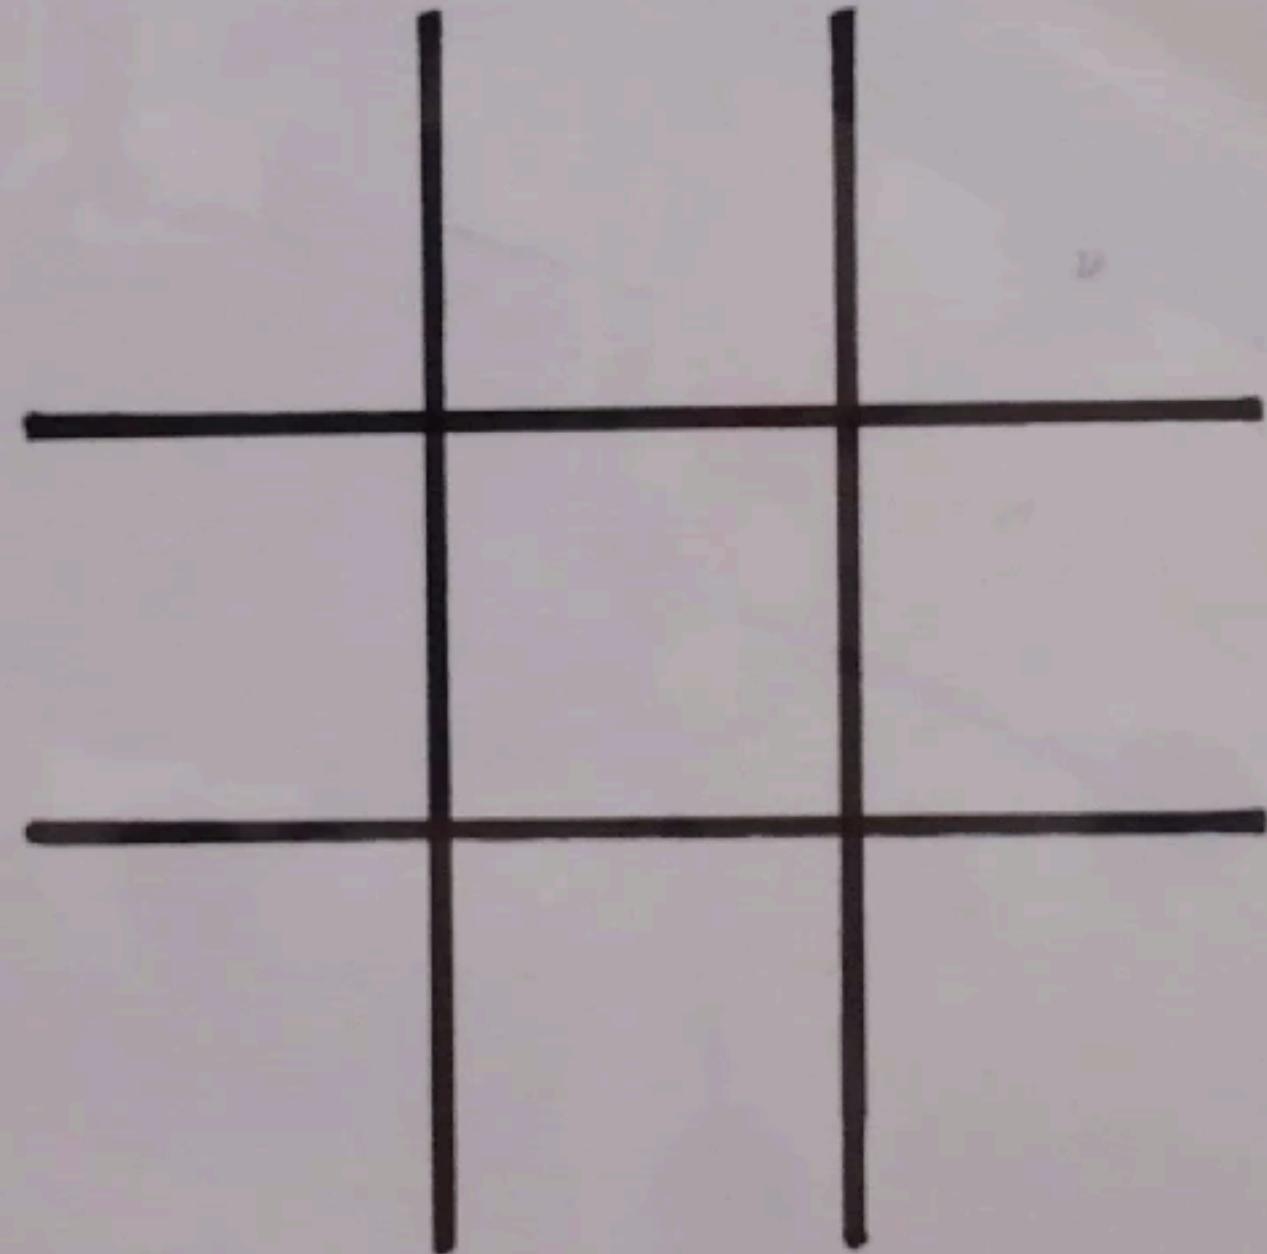
Ephemeral does not have a time limit.

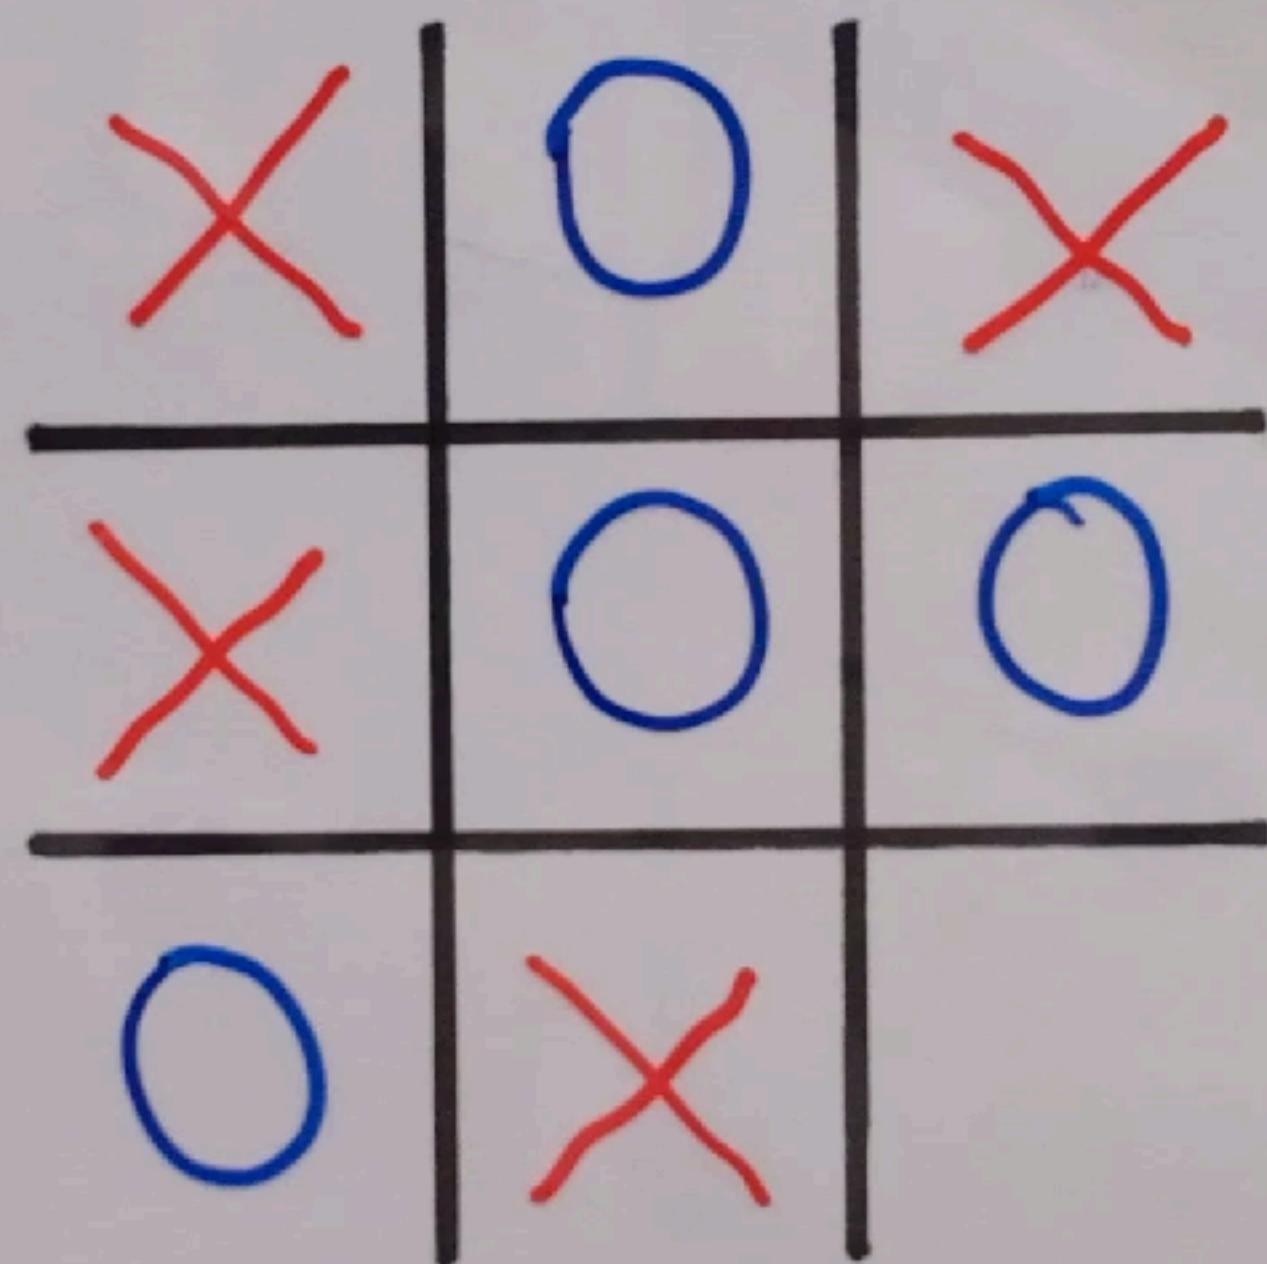
Ephemeral does not have a time limit.

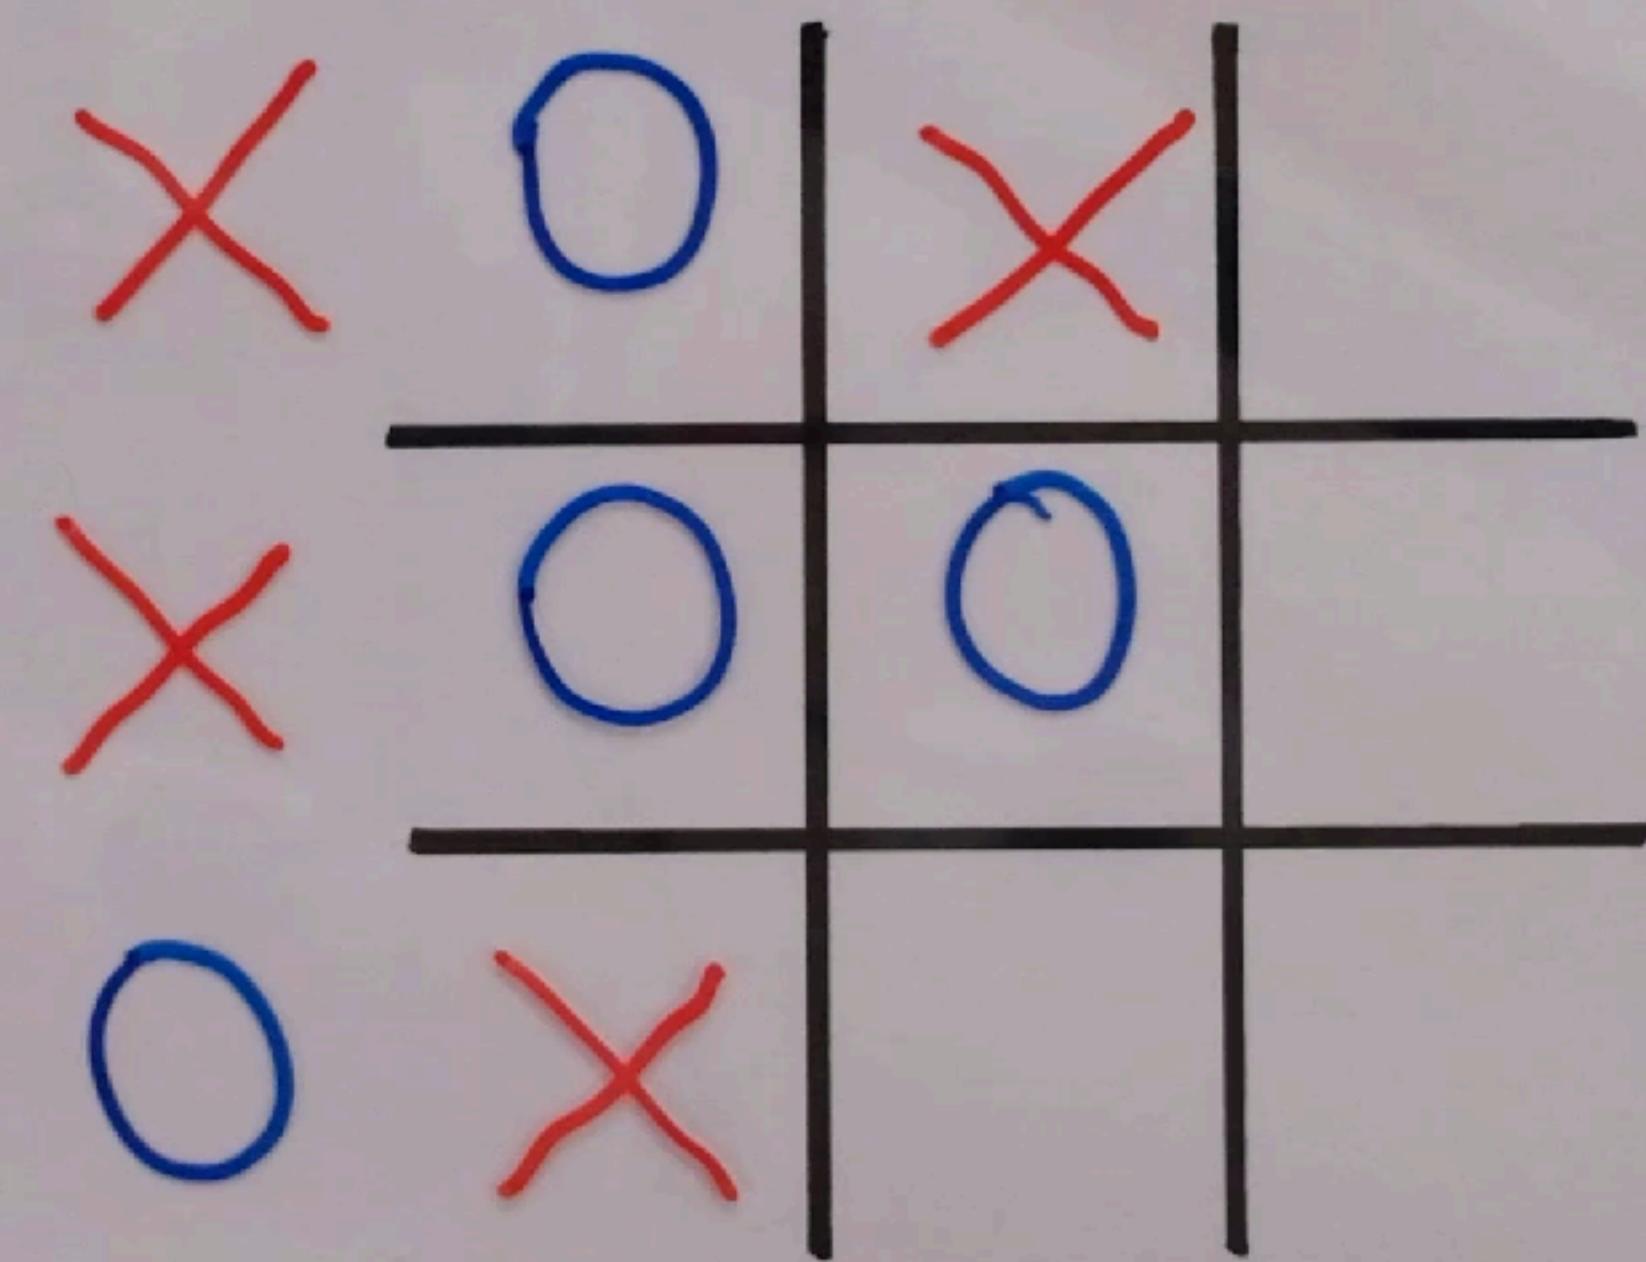
On a long enough scale,
everything is ephemeral.

Images are **immutable** and *never change*.

Stalemate!







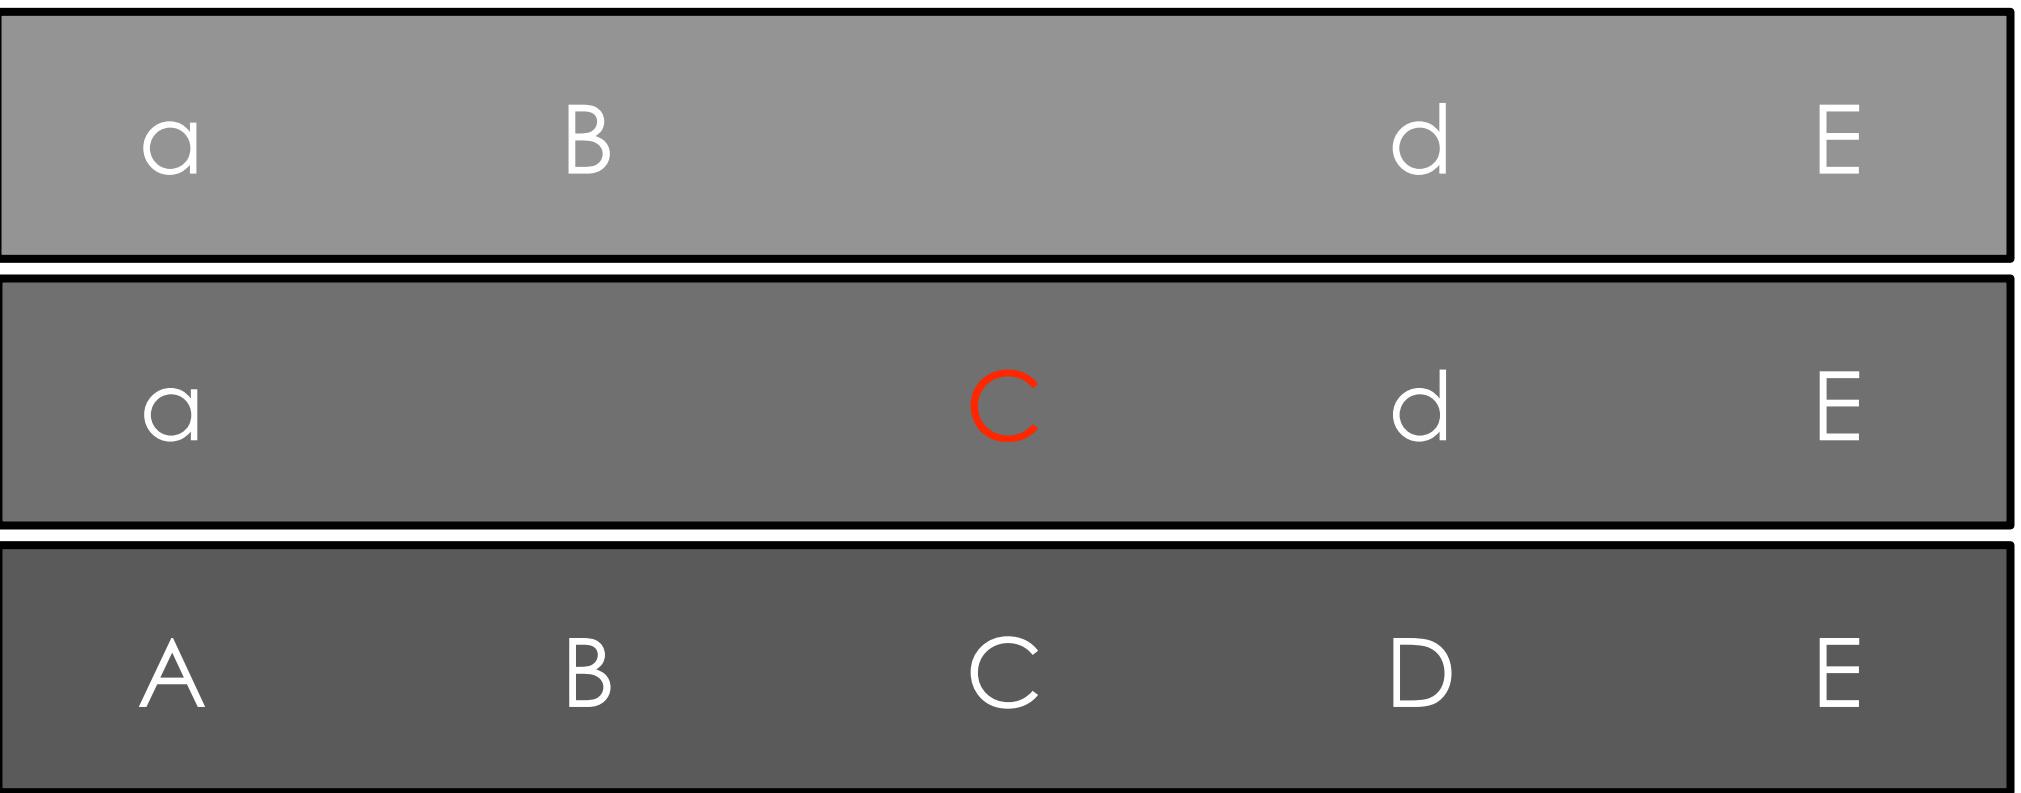
Volumes

Volumes "externalize" container directories.

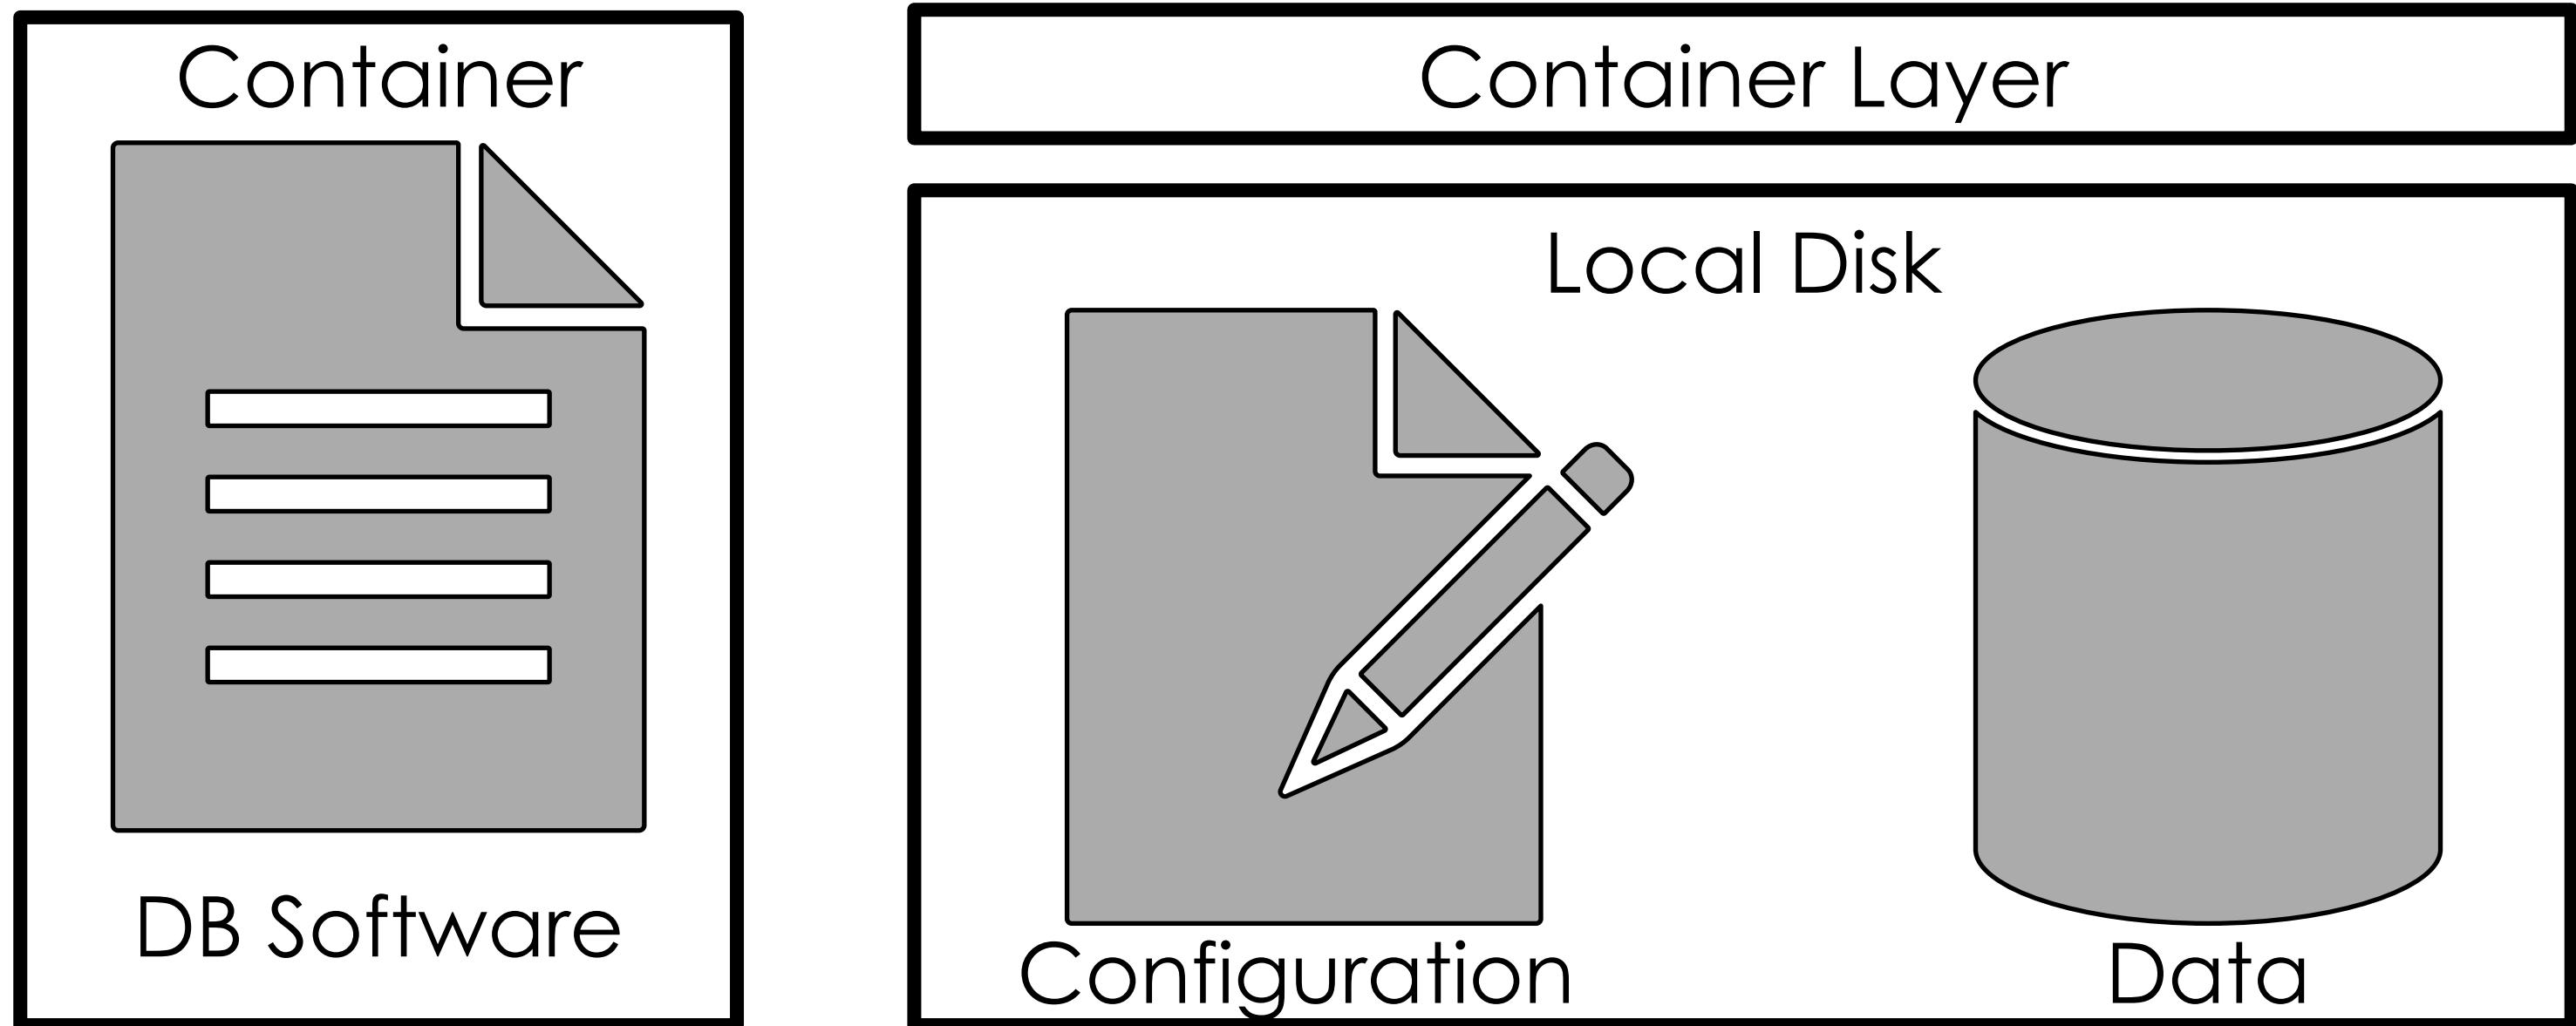
Union Layer

Container Layer

Image Layer



Database Containers



Database Image Rule #1:

If there is a database, start it.

If not, create a new database.

Database Container Startup

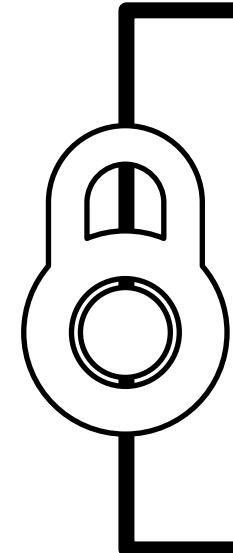
```
docker volume create ${CONTAINER_NAME}_data  
  
docker run -d \  
  --name ${CONTAINER_NAME} \  
  -v /data/${CONTAINER_NAME}:<data_directory> \  
  <image_name>
```

Default Volumes

Linux

```
ls -l /var/lib  
drwx--x--- root root docker
```

Docker Desktop



Virtual Machine

/var/lib/docker

```
mkdir -p /data/${CONTAINER_NAME}

docker volume create \
  --opt type=none \
  --opt o=bind \
  --opt device=/data/${CONTAINER_NAME} \
${CONTAINER_NAME}_data
```

Database Container Startup

```
docker run -d \  
  --name ${CONTAINER_NAME} \  
  -v /data/${CONTAINER_NAME}:<data_directory> \  
  <image_name>
```

Database Image Rule #1:

If there is a database, start it.

If not, create a new database.

Clone the Data Directory

```
cp -rp /data/${CONTAINER_NAME}/ \
    /data/clone
```

Clone the Data Directory

```
cp -rp /data/clone/ \
    /data/${CONTAINER_NAME}
```

Clone the Data Directory

```
cp -rp /data/clone/ \
    /data/${CONTAINER_NAME}
```

```
cp -rp /NFS/gold_data/clone/ \
    /data/${CONTAINER_NAME}
```

Clone the Data Directory

```
cp -rp /data/clone/ \
    /data/${CONTAINER_NAME}
```

```
cp -rp /NFS/gold_data/clone/ \
    /data/${CONTAINER_NAME}
```

```
cp -rp /NFS/gold_data/clone/2024-05-08.0900/ \
    /data/${CONTAINER_NAME}
```

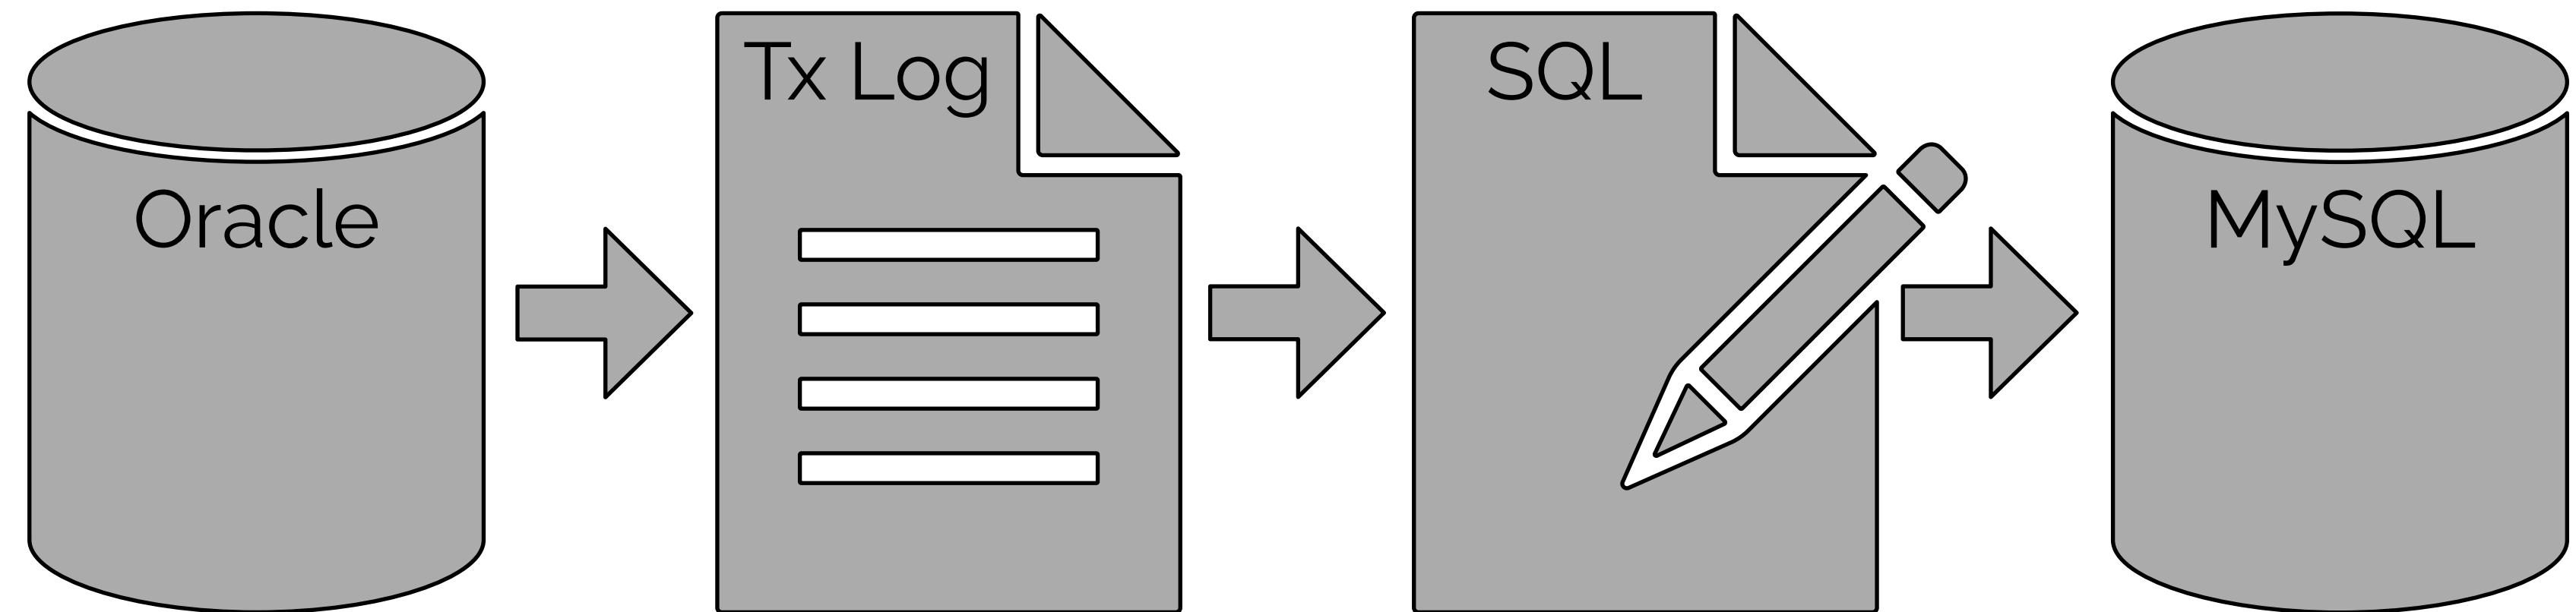
Data is Code

Data *is* Code

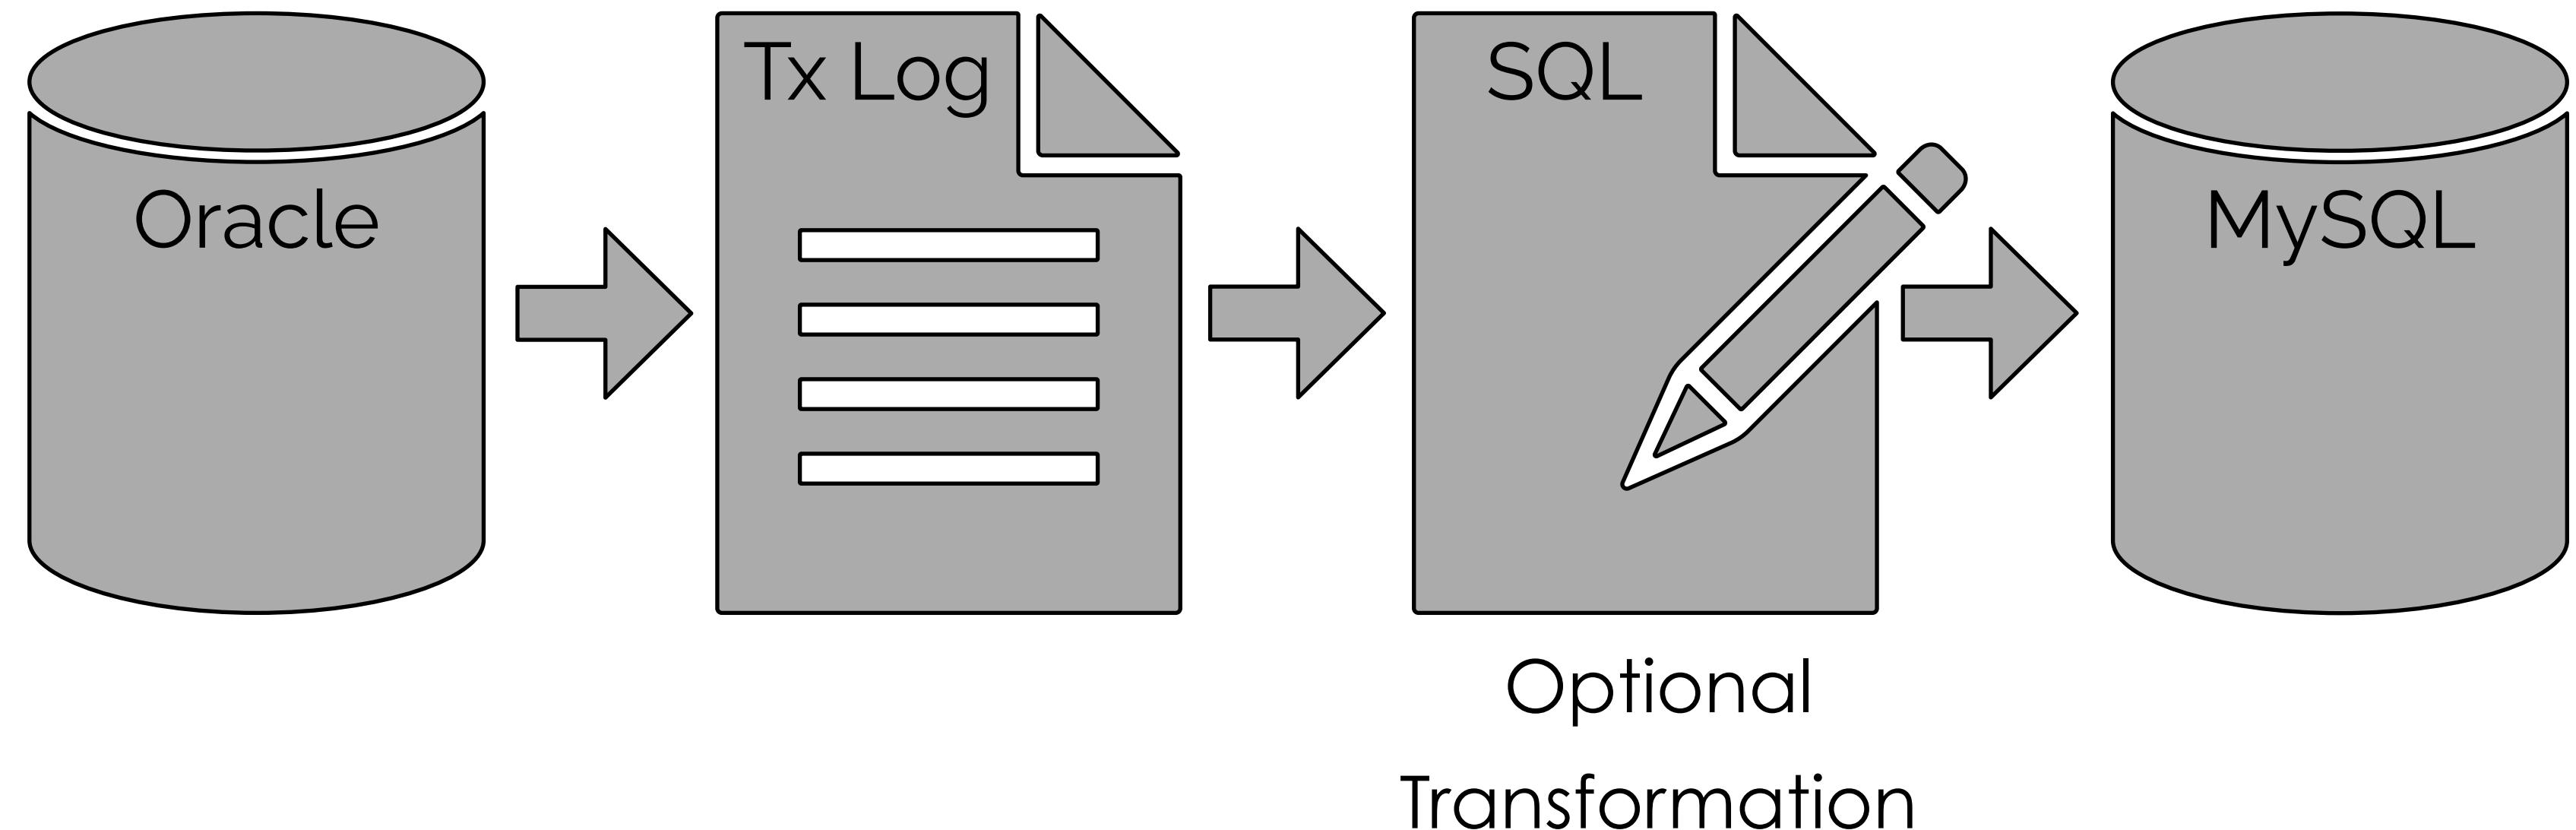
Every record in every database is created in one of three ways:

- INSERT
- UPDATE
- DELETE

Heterogenous replication = synchronous SQL



Heterogenous replication = synchronous SQL



If data can be codified, **data is versionable**.

Entrypoints

- Generally adapted but *not* an OCI standard
- Implementation varies by image, vendor/author
 - Setup only
 - Setup & startup

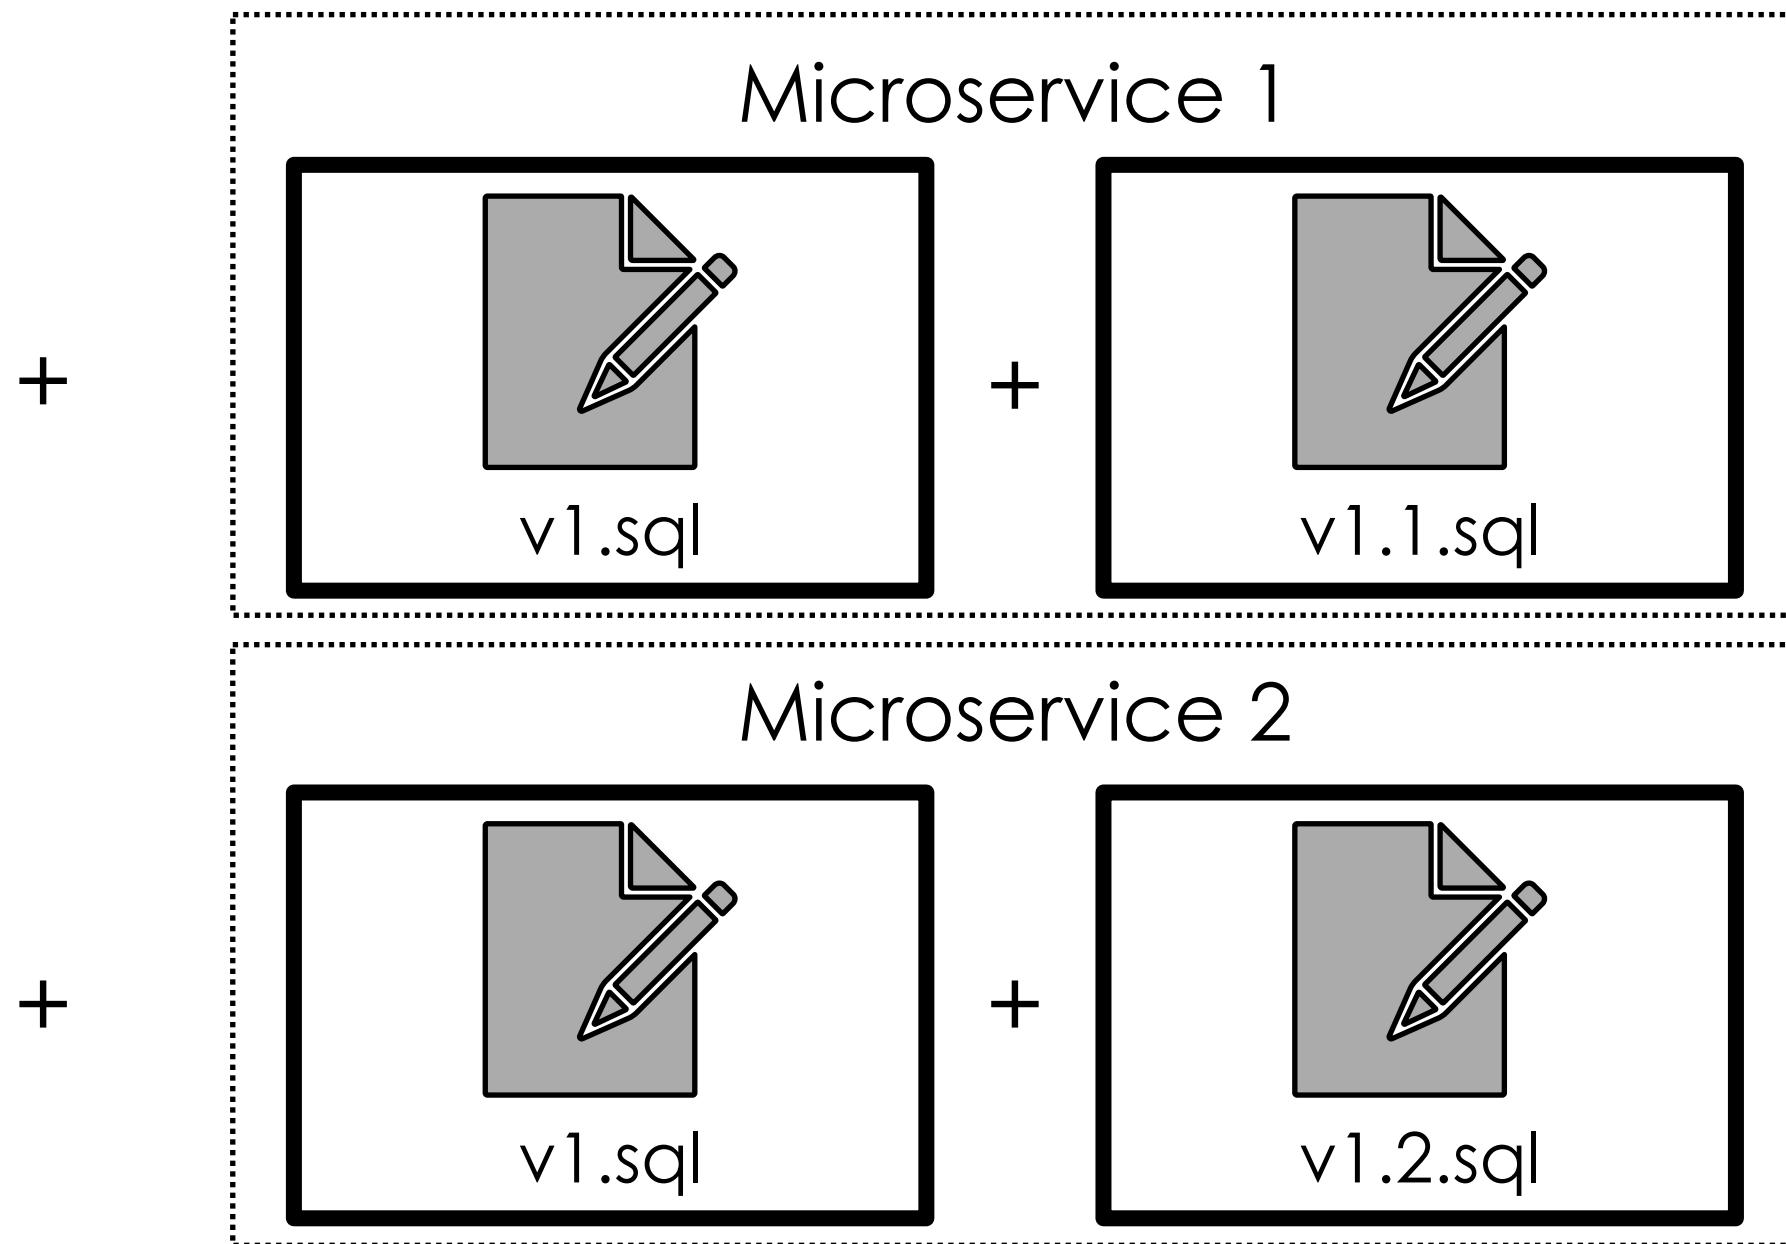
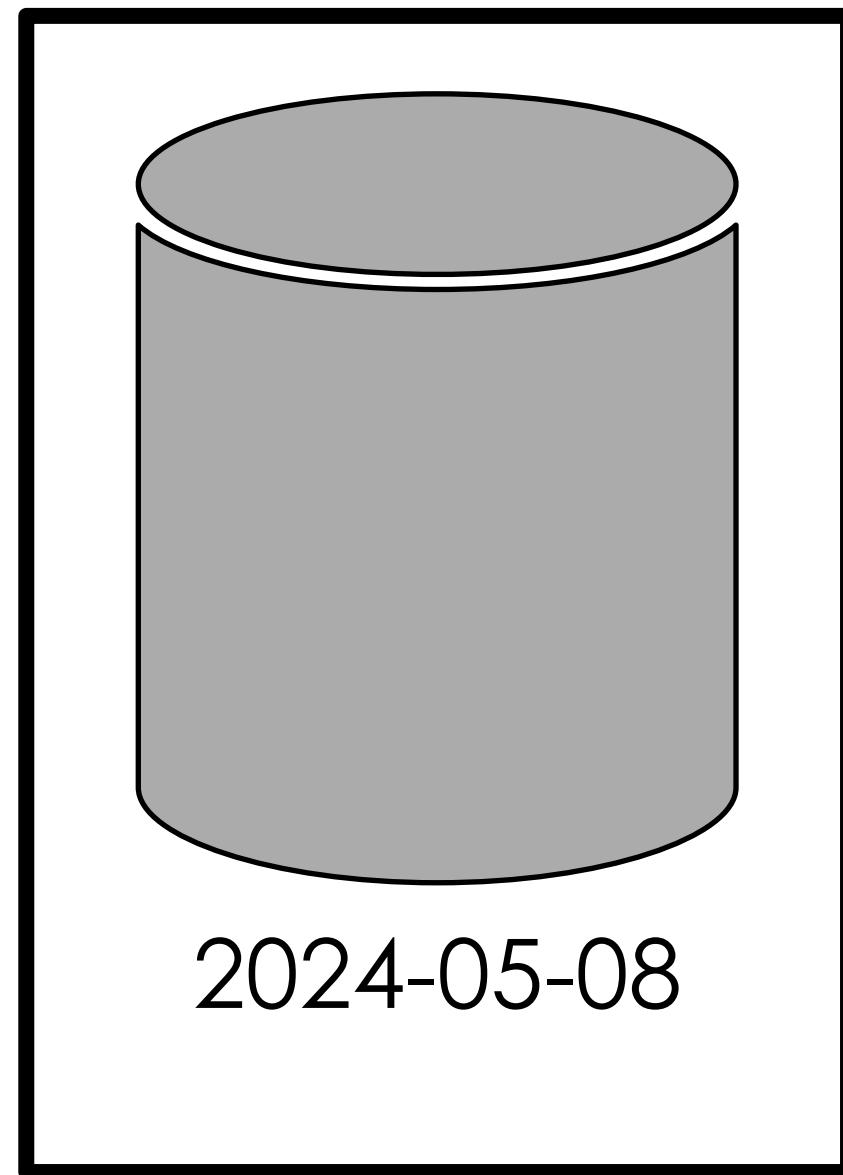
Entrypoints

```
COUNTAINER_NAME=V20240508
```

```
cp -r /data/2024-05-08/ \
    /data/${CONTAINER_NAME}
```

```
docker run -d \
    --name ${CONTAINER_NAME} \
    -v /data/${CONTAINER_NAME}:<data_directory> \
    -v /code/${CONTAINER_NAME}:/docker-entrypoint-initdb.d \
    <image_name>
```

Versioned Data



Building Images

A **Containerfile** is a recipe for building images.

Tic-Tac-Toe Grid

Paper

Image layers are *modular* and *reusable*.

Layers can be **cached** and reused
among multiple images.

Game of Sums

Very Tiny Sudoku

Tic-Tac-Toe Grid

Paper

```
# Create some directories:
```

```
mkdir -p ~/jotb/{image,c3_container,c3_work,container3}
```

```
# cd into the main directory:
```

```
cd ~/jotb
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=container1, \
upperdir=c3_container, \
workdir=c3_work \
overlay container3
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=container1, \
upperdir=c3_container, \
workdir=c3_work \
overlay container3
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=container1, \
upperdir=c3_container, \
workdir=c3_work \
overlay container3
```

```
# Update PATH (just in case!)
```

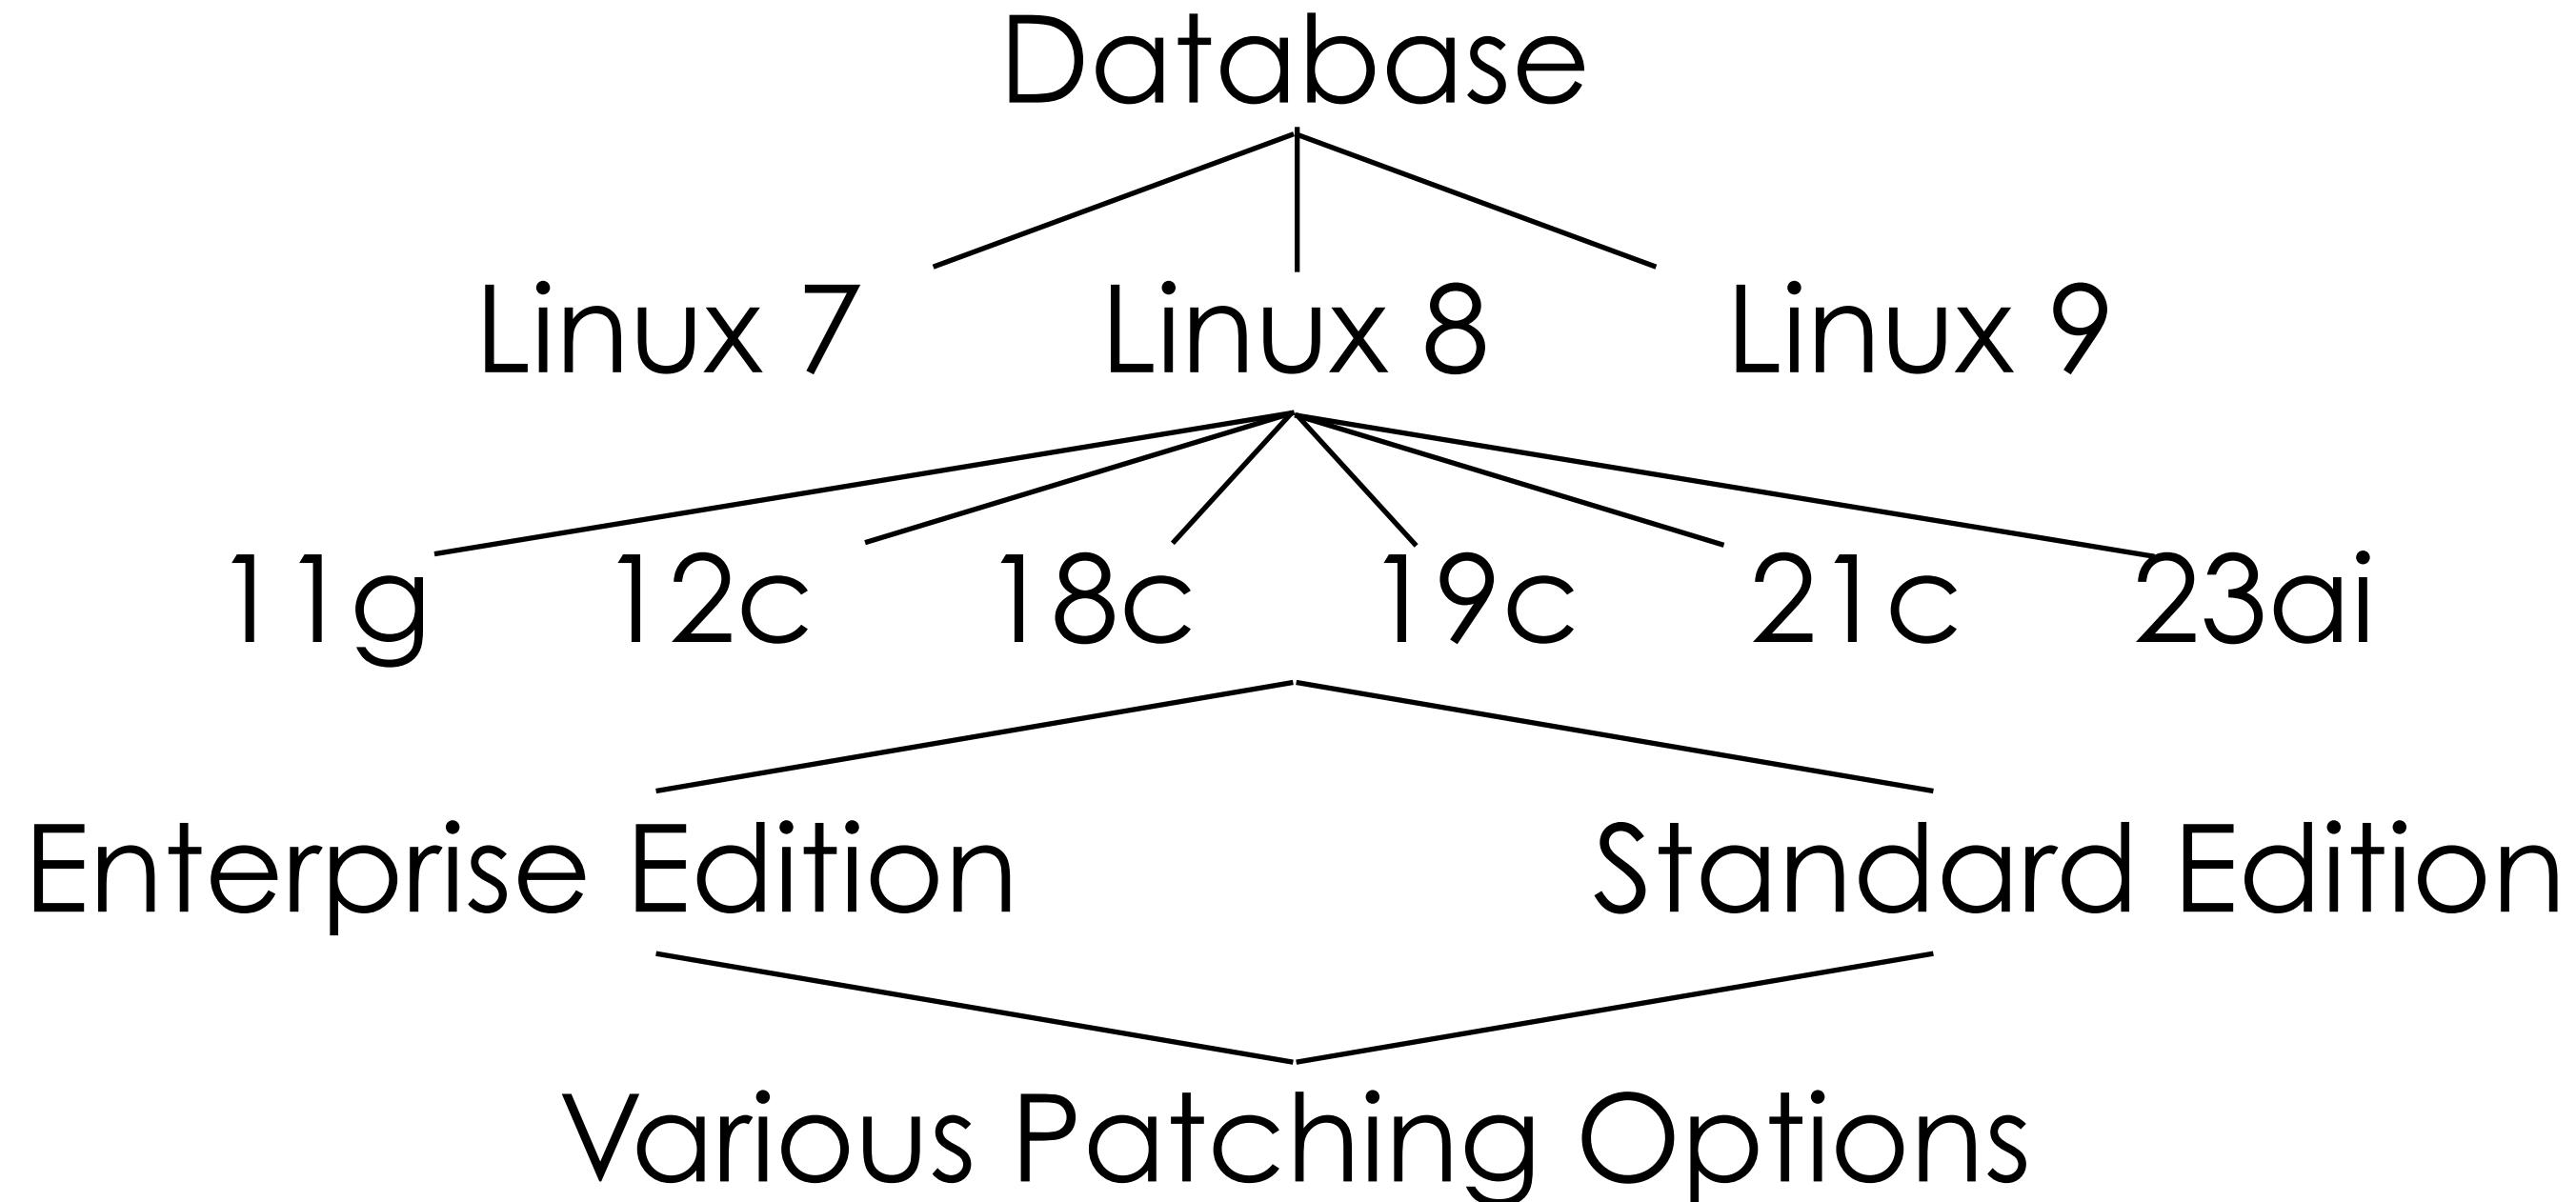
```
PATH=$PATH:/bin
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/jotb/container3 \  
/bin/sh
```

Context is a pantry of image ingredients.

Ignore files limit and control the build context.

Repository Structure



Questions

sean.scott@viscosityna.com
<https://linktr.ee/oraclesean>

www.viscosityna.com



Contact Me

sean.scott@viscosityna.com
<https://linktr.ee/oraclesean>

